



COLLEGE OF ENGINEERING, DESIGN AND PHYSICAL SCIENCES,
ELECTRONIC AND COMPUTER ENGINEERING

DISTRIBUTED COMPUTING SYSTEMS ENGINEERING

MASTER THESIS

Conception and Realization of a Distributed and
Automated Computer Vision Pipeline

Michael Watzko
1841795

Supervisor:
Dr. Paul Kyberd

March 31st, 2020

Abstract

A short summary of what the project is about.

TODO: .

schedule large work items

special hardware needs

automation and highly customizable pipeline

focus on easy setup and low maintenance

Keywords: Software, Architecture, Events, Messaging, Filesystem, Distributed, Coordination, Docker, Spring Boot, REST, Angular, Typescript

Contents

1	Introduction	1
1.1	MEC-View	3
2	Aims and Objectives	4
2.1	Current Workflow	4
2.2	Desired Workflow	6
2.3	Deliverable Requirements	7
2.3.1	Non-Requirements	8
3	Literature Survey	9
3.1	Similar solutions	9
3.1.1	Hadoop MapReduce	9
3.1.2	Build Pipelines	11
3.1.3	Camunda	12
3.1.4	Nomad	12
3.1.5	dCache	13
3.1.6	Further mentions	13
3.2	Docker	15
3.2.1	Technology	16
3.2.2	Ecosystem	17
3.3	Storage Technology TODO: Network File System	17
3.3.1	Remote File System	18

3.4	Angular Web-Application and REST API	19
3.5	yaml?	19
3.6	Agile development?	19
4	Introducing Winslow	20
4.1	Common Terminology	20
5	System Analysis	22
5.1	System Context	23
5.2	Use Case Diagrams	24
5.2.1	Managing Pipelines and Projects	24
5.2.2	Managing Resources and Workspaces	25
5.2.3	Managing and Monitoring Executions	26
5.2.4	Monitoring Nodes	26
5.2.5	System Administration	27
5.3	Detailed Use Case analysis	27
5.4	System Architecture	28
5.4.1	Layer 1: Orchestration Service	28
5.4.2	Layer 2: Events and Resources	29
5.4.3	Layer 3: Backend Driver	29
5.4.4	Layer 4: Client Communication	30
5.5	Time Schedule	31
6	System design	32
6.1	Environment	32
6.2	Storage Technology	32
6.3	Event Synchronization and Communication	35
6.3.1	Messages	37
6.4	Backend Driver	39
6.5	Execution Management	39
6.5.1	Remote Execution	39

6.5.2	Local Execution	40
6.5.3	The approach for Winslow	40
6.6	Backend Nomad	40
6.7	Large File Storage organization	41
6.7.1	Stage storage	41
6.8	Job Scheduling / Election System	43
6.9	CPU RAM Network IO Monitoring	43
6.10	User Interface	43
6.11	File system	44
6.11.1	Every node has a connection to every other node	44
6.11.2	Centralized broker	44
6.11.3	Tree hierarchy	44
6.11.4	outcome	44
6.12	Communication and Node Management	44
6.12.1	Centralized Management, Remote Execution	44
6.12.2	Decentralized Management, Local Execution	44
6.12.3	Combinations worth noting	44
6.12.4	Architecture	45
6.12.5	Failure handling	45
6.13	Communication/Event architecture?	46
6.14	Targeting capabilities	46
6.14.1	General thoughts	46
6.15	Planned	46
6.16	Implementation details	46
6.17	Synchronization, Locking, publishing events	46
6.18	REST for UI	46
6.18.1	Atomicity of (Unix) Filesystems	46
6.18.2	Atomicity and behavior of NFS in particular	46
6.18.3	Using as lock backend	46

6.18.4 Using as election backend	46
6.19 Agile development	46
6.20 Continuous Deployment	46
7 Outcome and Measurements	48
7.1 What did not work as expected	48
7.2 Storage	48
7.3 Coordination	49
7.4 Binary distribution	49
7.5 User Interface	50
7.6 Requirements	50
7.7 Architecture	50
7.8 High Level System Overview	50
7.9 Evaluation	52
8 Further work	56
9 Conclusion	57
Bibliography	61

Chapter 1

Introduction

Since the industrial revolution, humans strive for more automation in the industry as well as in the every day life. What was at first a cost saving measurement in factories, now also is a differentiation method for products. A new product must prove a higher level comfort to the customer than the previous generation as well as all the competitors. As such, the ambitions of the industry are focused on increasing the value of their products for the customer.

The automotive industry is one of the prime examples of this. Never was traveling from one place to another as comfortable as nowadays. Aspects like an elegant interior design, comfortable seats, air conditioning, entertainment systems and safety measurements need to be considered by car manufacturers to be competitive these days. The next luxury enhancement will be the autonomously driving vehicle. No longer shall the owner of a car steer it, but instead the car becomes his or hers personal chauffeur, driving the optimal route, the most comfortable way and being more reliable and safer than any human ever could.

The reason, autonomously driving cars are not common already, is their big complexity increase. Compared to already established technologies like parking assistants, entertainment systems or more efficient engine controllers, letting a computer reliably understand a certain traffic situation requires masses of input data and complex algorithms to process. As such, the problem itself becomes

massive and cannot be solved that easily. So the industry has no choice than to divide this into many small pieces and work out solutions step by step.

The MEC-View research project explores one such step: whether and how to include external, steady mounted sensors in the decision finding process for partially autonomous vehicles in situations where onboard sensors are insufficient. To not disrupt traffic flow with non-human behavior, one needs to study and thereby watch human traffic. Automatically analyzing traffic from video footage requires a lot of computation power and can be further optimized by specialized hardware such as GPUs¹.

This thesis will conceptualize and realize a distributed and automated computer vision pipeline which is in this case used to analyze traffic flow within video footage. Compared to an existing but highly manual workflow, the new system shall help to utilize the available hardware more efficiently by reducing idle times. Stage transitions and basic scheduling shall be automated to allow a user to plan and execute multiple projects ahead of time and in parallel.

¹Graphics Processing Units

1.1 MEC-View

TODO: shorten!?

The MEC-View research project[1] - funded by the German Federal Ministry for Economic Affairs and Energy - aims to supplement the field of view of automated driving cars with road-side sensor data using 5G mobile communication. The sensor information is merged into an environment model on the so-called Mobile Edge Computing (MEC) server. This server is directly attached to the radio station to ensure low latency environment model updates.

The project is tested at an intersection in Ulm, Germany. Currently, there are 15 lidar and video sensors installed. Those sensors send their detections to the (MEC) server. A fusion-algorithm merges those detections into one environment model and sends it back to the (MEC) server and to the automated cars.

Additionally, general traffic flow is analyzed to learn about movement patterns. To do so, 4k video data is captured by an air drone from real world cross roads - not limited to the intersection in Ulm. On each frame of such a recording, cars are detected with a neuronal network. Detected cars are tracked throughout the video to compute the movement speed and position in time of each car. In an analysis of all vehicles, hot-spots of high and low traffic flow can be determined.

This thesis conceptualizes and implements a distributed and automated computer vision pipeline to increase the efficiency in video analysis management and execution. It is of concern for this thesis on how to retrieve the footage or what is further archived with the results of evaluated video footage. The focus is on utilizing available hardware resources to manage multiple projects simultaneously and to reduce the idle time of relevant hardware by slow human reaction and availability times. TODO: reaction time: like its done but not noticed for another x mins/hours, availability: work hours

Chapter 2

Aims and Objectives

This chapter will discuss the program which shall be implemented. To do so, the problem to solve must be understood. To gather requirements and understand the technical hurdles to overcome, this chapter is split into two sections. First, a rough glance over the current workflow is given, which is followed by a more detailed description for the desired workflow.

2.1 Current Workflow

Currently, to analyze a video for the trajectories of recorded vehicles, the following steps are executed manually:

1. Upload the input video to a new directory on the GPU server
2. Execute a shell script with the video as input file and let it run (hours to days) until completed. The shell script invokes a Java Program - called TrackerApplication - with parameters on what to do with the input file and additional parameters.
3. The intermediate result with raw detection results is downloaded to the local machine and opened for inspection. If the detection error is too high,

the camera tracking has a drift or other disruptions are visible, the previous step is redone with adjusted parameters.

4. Upload the video and intermediate result to a generic computing server and run data cleanup and analysis. This is achieved with the same Java Program as in step 2, but with different stage environment parameters.
5. Download the results, recheck for consistency or obvious abnormalities. Depending on the result, redo step 2 or 4 with adjusted parameters again.
6. Depending on the assignment, steps 4 and 5 are repeated to incrementally accumulate all output data (such as statistics, diagrams and so on).

Because all those steps are done manually, the user needs to check for errors by oneself. Also, if a execution is finished or has failed early, there could be hours wasted until noticed, if the check intervals are too far apart, such as during nights or weekends.




Figure 2.1: Overview of workflow

2.2 Desired Workflow

The desired workflow shall be supported through an user interface that provides an overview of all active projects and their current state, such as running computation, awaiting user input, failed or succeeded.

To create a new project, a predefined pipeline definition shall be selected as well as a name chosen. Because only a handful of different pipeline definitions are expected, the creation of such does not need to happen through the user interface. Instead, it is acceptable to have to manually edit a configuration file in such rare circumstances.

Once a project is created, the user wants to select the path to the input video. This file has to be been uploaded to a global resource pool at this point. The upload and download of files shall therefore also be possible through the user interface. Because a video is usually recorded in 4k (3840 x 2160 pixels), encoded with H.264 and up to 20 minutes long, the upload must be capable of handling files which are tens of gigabytes large.

Once a pipeline is started, it shall execute the stages on the most fitting server node until finished, failed or a user input is required. Throughout, the logs of the current and previous stage shall be accessible as well as uploading or downloading files from the current or previous stages workspace. In addition to the pipeline pausing itself for user input, the user shall be able to request the pipeline to pause after the current stage at any moment. When resuming the pipeline, the user might want to overwrite the starting point to, for example, redo the latest stage.

Mechanisms for fault tolerance shall detect unexpected program errors or failures of server nodes. Server nodes shall be easily installed and added to the existing network of server nodes. Each server node might provide additional hardware (such as GPUs), which shall be detected and provided.

For the ease of installation and binary distribution, Docker Images shall be used for running the Java Program for analyzing the videos as well the to be implemented management software.

2.3 Deliverable Requirements

From the desired workflow, the following requirements can be extracted **TODO:** (shortened and incomplete due to early project stage):

- User interface for interaction between the system and the user
- Storage management for global resource files as well as stage based workspaces
- Pipeline definition through configuration files
- Handling of multiple projects with independent progress and environment
- Reflecting the correct project state (running, failed, succeeded, paused)
- Log accumulation and archiving
- Accepting user input to update environment variables, resuming and pausing projects as well as uploading and downloading files into or from the global resource pool or a stages workspace.
- Assigning starting stages to the most fitting server node
- Detecting program errors (in a stage execution)
- Cope with server node failures
- **TODO: CI** Docker Image creation for the Java Binary as well as the program implementation, preferred in an automated fashion.

Furthermore, a given constraint is that the user interface is to be implemented as Angular Web-Application using a REST API (for more details see ??).

TODO: something something agile extended as needed? section 3.6

2.3.1 Non-Requirements

To know the requirements and expectations of a system is essential, but knowing what is not expected by the system is at least as valuable. It prevents wasting resources, effort and architectural specialization that will never be required or in the worst case, make further development harder by restricting available choices in the future.

The system to implement shall not strive to implement real time scheduling or low latency scheduling. The expected work items are big chunks that require hours to compute. Whether the assignment of the work item takes sub seconds or several seconds is nearly unnoticeable in the overall compute time.

Chapter 3

Literature Survey

This chapter is about accumulating information. Programs that are solving similar problems, as described in the desired workflow, or dealing with a subset of the problem are looked into. Fundamental knowledge for understanding this thesis is also acquired here.

3.1 Similar solutions

This sections focuses on programs trying to provide somewhat similar workflows. The reason for this is to use well established or suitable programs as middle-ware to reduce implementation overhead, or, where this is not possible, one might be able to gather ideas and learn about proven strategies to use or pitfalls to avoid while implementing custom solutions.

3.1.1 Hadoop MapReduce

For big data transformation, Hadoop MapReduce[3] is well known. With MapReduce, the input data is split into blocks of data and distributed onto Mappers. Mappers then apply the business logic and output intermediate results in form of key/value pairs. After shuffling, the Reduce stage will combine values from key/-value pairs with the same key. In the final output, each key is unique accompanied

with a value.

This strategy has proven to be very powerful to process large amount input data because Mappers and Reducers can work independently on their data-sets and therefore scale very well when adding further instances.

If the implementation were to be based on Hadoop MapReduce to achieve the desired workflow, it could be done like the following:

- Each video is split into many frames and each frame is applied to a Mapper
- A Mapper tries to detect all vehicles on a frame and outputs their position, orientation, size and so on
- The Reducer then tries to link the detections of a vehicle through multiple frames
- The final result would be a set of detections and therefore all positions for each vehicle in the video

But at the moment, this approach seems to be unfitting due to at least the following reasons:

1. It is not always trivial to reasonable link the detections of a vehicle. For example, a vehicle can be hidden behind a tree for a few frames until visible again. In addition, MapReduce requires the combination to be performed per common key. Until one is trying to link the detections of multiple frames, there is no common identifier that could easily be used as key. The position of a moving vehicle cannot be used as key, neither can the color or size, because of the noise of the camera, deviation in detection output and perspective distortions. The current implementation of the TrackerApplication is archiving this by finding similarities between detections, but for the Mapper it would be required to express this as a deterministic key.
2. MapReduce is great in combining many machines to solve a big computational problem. But at the moment, this is neither a desired nor given

condition. At the moment, there are a handful of very powerful workstations with specialized hardware. Therefore it is perfectly acceptable and sometimes required, when each workstation works through a complete video at a given time instead.

3.1.2 Build Pipelines

Build pipelines such as GitLab[4] and Jenkins[5] can also distribute the execution of stages onto other server nodes. In a common use-case, such build pipelines are used to build binaries out of source code, after a new commit into a SCM¹ repository was made. At IT-Designers GmbH GitLab as well as Jenkins are commonly used for scenarios exactly like this. A pipeline definition in GitLab CI/CD [6] or in a Jenkinsfile [7] describe stages and commands to execute. Each stage can be hosted on another node and be executed sequential or in parallel to each other.

Although this seems to be quite fitting for the desired workflow, there are two issues. First of all, such a pipeline does not involve any user input besides an optional manual start invocation. The result is then determined based on the state of the input repository. Second, such a pipeline is designed to determine the output (usually by compiling) whereas each run is independent from the previous and a repeated run shall provide the same result as the previous did. Usually, a new run is only caused by a change of the input data. However, the desired workflow differs in this aspects. A redo of a stage can depend on the result of the previous stage, for example, if the results are poor or the the stage failed. Instead of having multiple complete pipeline runs per project, the desired workflow uses a pipeline definition as base for which the order can be changed. Also, intermediate results need to influence further stages, even if repeated.

¹Source Code Management

3.1.3 Camunda

Camunda[8] calls itself a “Rich Business Process Management tool” and allows the user to easily create new pipelines by combining existing tasks with many triggers and custom transitions. Camunda is focused upon visualizing the flow and tracking the data through a pipeline. The Camundas Process Engine[9] also allows user intervention between tasks.

One of the main supporting reason for it Camunda is the out of the box rich graphical user interface for process definition and interaction. Through its API[10], Camunda also allows custom external workers to execute a task. But it misses the capability to control which task shall be processed on which worker node which is required by the desired workflow. It does also not provide any concept on how to allocate and distribute resources. The user interface - while being rich overall - is quite rudimentary when it is about configuring tasks and would therefore require custom plugins to be developed for more advanced user interactions.

Camunda is also not designed to reorder stages or insert user interactions at seemingly random fashion. The user itself is considered more as a worker that gets some request, “executes” this externally and finally marks the request as accepted or declined. Mapping this to the desired workflow does not feel intuitive. Finally, there is also no overview of task executors, no centralized log accumulation and no file up- or download for global project resources.

3.1.4 Nomad

Nomad[11] by HashiCorp is a tool to deploy, manage and monitor containers, whereas each job is executed in its own container. It provides a rich REST API and can consider hardware constraints on job submissions. Compared to Kubernetes[12], which is similar but more focused on scaling containers to an externally applied load, it is very lightweight. It is also available in many Linux software repositories - such as for Debian - which makes the installation very easy.

Because there were no grave disadvantages found (depending on a third party library can always be considered be a disadvantage for flexibility, error-pronous and limit functionality) Nomad is being considered as a middle-ware to manage and deploy stages. Others[13] seem to be using Nomad to manage and deploy containers for similar reasons. Nonetheless, further testing and prototyping will be required for a final decision.

3.1.5 dCache

“The goal of this project is to provide a system for storing and retrieving huge amounts of data, distributed among a large number of heterogenous server nodes, under a single virtual filesystem tree with a variety of standard access methods”[14]. dCache seems to be able to solve the storage access and distribution concern for the stages and sever nodes. When using dCache, one could store the global resources distributed between the server nodes. Built-in replication would prevent access loss on a node or network failure and an export through NFS² allows easy access for Linux based systems[15].

But the installation is complex and requires many services to be setup correctly, such as postgresql and many internal services such as zookeeper, admin, poolmanager, spacemanager, pnfsmanager, cleaner, gplazma, pinmanager, topo, info and nfs. The documentation is also rather outdated and incomplete which meant, early tests with a prototype setup took days to setup and behaved rather unstable (probably due to a wrong configuration). It is to be seen, whether such an complex and heavy system is actually required or if there are feasible alternatives.

3.1.6 Further mentions

The following list acknowledges programs that behave similar to the previously mentioned strategies or can be used as building blocks for custom implementa-

²Network File System

tions. Programs that are listed here, were looked into, but not all in great depth because miss-fits or missing functionality were detected early on. The list is in no specific order:

- **Quartz**[16] is a Java based program to schedule jobs. Instead of doing so by using input, Quartz executes programs through a timetable and in certain intervals.
- **Luigi**[17] also executes pipelines with stages and is written in python. The advertised advantage is to define the pipeline directly in python code. But, this is at the same time the only way to define pipelines which contradicts with the existing Java TrackerApplication implementation.
- **Calery**[18] is focused on task execution through message passing and is written in Python. Intermediate results are expected to be transmitted through messages. Because there is no storage strategy and python adapter-code would have been required, Calery was dismissed.
- **IBM InfoSphere**[19] provides similar to Camunda a rich graphical user interface but for data transfer. Dismissed due to commercial nature.
- **qsub**[20][21] is a CLI³ used in HPC to submit jobs onto a cluster or grid. Dismissed due to an expected high setup overhead, non-required multi-user nature and the fact, that it only provides a way to submit jobs.
- **CSCS**[22] High Throughput Scheduler (GREASY). Dismissed for similar reasons as qsub, although it is more light weight and hardware agnostic (it can consider CUDA/GPU requirements).
- **zsync**[23], similar to rsync, is a file transfer program. Zsync allows to only transfer new parts when a file that shall be copied already exists in an older version on the target. This tool might be useful when implementing a custom resource distribution strategy is required.

³Command Line Interface

- **OpenIO**[24] provides a distributed file system, is already provided as Docker image and provides a simple to use CLI. Because the NFS export is only available through a paid subscription plan, it was dismissed from further investigation.
- **SeaweedFS**[25] provides a scalable and distributed file system. The most interesting aspects are that it is rack-aware as well as natively supports external storage such as Amazon S3. When adding server nodes from the cloud this could allow all nodes to access the same file system while using rack-aware replication to reduce bandwidth usage and latency. A local test also proved that it is easy to setup, but because it cannot hot-swap nodes and was not able to recover when the seaweed master node became unreachable it was dismissed.
- **Alluxio**[26] provides a distributed file system but was dismissed because it itself requires a centralized file system for the master and its fallback instances
- **GlusterFS**[27] is another tool to provide a distributed file system with replication. It was bought by IBM but is nonetheless available through the software repository of many Linux distributions such as Debian. A local test showed that the setup is very easy and no adjustments of configuration files are required. However, the replication mechanism requires that an integer multiple of nodes of the replica value are assigned to the file system. This makes GlusterFS hard to use in a scenario, where adding and removing nodes are expected to happen frequently.

3.2 Docker

As describe in section 2.2, for easy deployment, the implementation as well as the stages shall be executed inside Docker[28] containers. This allows easier isolation

of the stages and workspaces from each other and other host programs, as will be explain further on.

Docker is the name of a software that combines isolation technologies (subsection 3.2.1) and a rich ecosystem (subsection 3.2.2) to provide and execute third party software in virtual environments. Docker aims are to increase security and to simplify installation as well as maintenance of applications. Docker Swarm[30] and Kubernetes[31] use these fundamental features to scale applications, services and microservices in the cloud.




Figure 3.1: Official Docker “Moby” Logo[29]

3.2.1 Technology

TODO: .

Docker uses so called images to package and transport binaries with all their required libraries and configuration files as a read-only archive to the destination system. Instead of spawning a new process for a binary in the host environment, the image is used to create a new virtual environment - the so called container. Because the image includes all required libraries, it is possible to completely isolate the virtual environment which makes the host system invisible to any process inside the container. Changes to files within containers are also stored separately⁴, which allows an image to be used by multiple containers at once. Privileges, resource limitations and storage configurations can also be specified when creating a new container. Processes inside containers are unable to see other processes or files that are not part of or assigned to the their container⁵.

In contrary to a hypervisor, docker is archiving this without hardware and operating system virtualization. Instead of executing the container inside an

⁴**TODO: overlayfs, similar to differentials in backups**

⁵This is the default behavior. It is possible to manually lift or modify many boundaries Docker enforces for containers on default.

additional virtualized operating system, they share the kernel of the host system. For that, the host needs to support additional isolation mechanisms. At the time of writing this, only the Linux Kernel is capable to separate processes, network interfaces, interprocess communications, filesystem mounts and the timesharing systems by namespaces. By configuring these namespaces, Docker is capable to isolate containers into virtual environments. Furthermore, control groups can be used to limit and constraint access to hardware resources. [32]

These approaches allow containers to run with very little additional overhead in comparison to running the application directly on the host. Containerization increases security by limiting what an application sees and is able to interact with, decreases maintenance overhead because of no additional operation systems to maintain and allows to run multiple instances of the same application besides each other with independent configurations and environments but with the same base image.

3.2.2 Ecosystem

TODO: Comment: Further Docker topics

3.3 Storage Technology **TODO: Network File System**

When distributing workload onto different machines, accessing input and output files becomes another concern which is not present when computation and storage reside on the same physical machine. In theory, one could use portable mediums such as USB-Sticks, CD-ROMs or external HDDs, but in reality, this becomes very tedious very fast. More advanced users might be able to take advantage of a common network connection between the computation and storage nodes to copy files and directories from and to the required places. This strategy -

still being tedious - surfaces another issue: dealing with multiple copies requires careful version management and additional storage space. One certainly would not want to continue computation on outdated data or have too many copies of the same files.

The solution is to use what is called network file system. **TODO: welp.** To programs they seem to be just another local directory hierarchy, but in reality, the files might be located on another or multiple other machines.

TODO: Comment: POSIX Stream API

3.3.1 Remote File System

Remote file systems provide access to data that does not need to be stored on the local machine and can often be shared with other clients as well. Some file systems are organized centralized, where at a single address and physical location all data is stored and retrieved from. This can have advantages in setup time, maintenance and communication complexity, because of their simplicity.

The Network File Systems (NFS) for example, which is primarily used in the Linux and Unix world, is centralized, the clients are connecting to a remote server to read and write files from and to. Since NFS 4.1 pNFS⁶[33, p. 14, section 1.7.2.2] allows the server to spread load across multiple servers and volumes, but because all metadata is still handled by one server, this approach still has a single point of failure. The Common Internet File System (SMB/CIFS) is similar to NFS in regards of providing access to a remote file system locally, and it is the approach primarily used by Windows computers. Similar to NFS, it has also a centralized approach.

More advanced file systems provide additional functionality like replication, which stores the data more than once to recover from data carrier failures, site awareness, which replicates data to geographically distant locations⁷ and decen-

⁶parallel NFS

⁷to not be affected by a burning data centre for example or to decrease access time for clients from the distant location

tralized communication endpoints that provide access to the data from more than a single point of failure.

File systems like GlusterFS (subsection 3.1.6), Alluxio (subsection 3.1.6), SeaweedFS (subsection 3.1.6), OpenIO (subsection 3.1.6), dCache (subsection 3.1.5) and HadoopFS (subsection 3.1.1) try to provide these functionality.

TODO: .

3.4 Angular Web-Application and REST API

TODO: HTTP

TODO: REST

TODO: Typescript

TODO: Angular, very structures, helps in development

TODO: upload/download -> difference HTTP Header/Request and Body?

3.5 yaml?

3.6 Agile development?

Chapter 4

Introducing Winslow

The program that shall implement the requested features listed by ?? is being called Winslow. This name refers to Frederick Winslow Taylor who was an American mechanical engineer and one of the first management consultants in the 19th century[.] Both strive to make people's work more efficient.




Figure 4.1: Winslow Logo

4.1 Common Terminology

This chapter explains the meaning of words and terminology used later on. It is crucial that every reader has the same understanding for the words used so that there are no wrong assumptions, expectations or surprises.

The root of a work item is called project. A project will usually refer to one video footage that shall be processed. Each project has its exclusive directory for input, output and intermediate files, called workspace. Furthermore a project as an author and possibly participants, that can help in steering and monitoring the processing. To do so, there is always a pipeline assigned to a project.

A pipeline consists of at least one but usually multiple stages that can at least be processed linearly - one after the other. Furthermore, a pipeline can define environment variables that valid for all stages within the stage.

A stage is the smallest work unit that can be executed. A docker image section 3.2 is specified as execution environment as well as further stage specific environment variables, command line arguments and hardware requirements (such as CPU cores, GPUs and minimum available RAM). This enables a stage to use a common image as well as to specify very precisely how to process the data.

In addition, stages and pipelines can be distinguished in a active, running or completed stage or pipeline and a definition. A stage definition specifies the above mentioned presets - allowing the user to adjust details just before submission - to create multiple stages from, while a running or complete stage has all information of what is or was executed and allows no further alteration. For pipeline definitions this is similar, a pipeline definition can be used to specify a common order of execution. Once assigned to a project, the project can freely adjust and change its instance of the pipeline.

A running process that is executing the binary of Winslow is called Winslow instance.

A computer with a Winslow instance that is accepting or executing stages is called execution node. The focus relies here on the computer providing compute resources to Winslow.

TODO: cleanup

Chapter 5

System Analysis

This chapter is about analyzing the to be implemented system more detailed, separating concerns **TODO: cite?** and starting to evaluate solutions in an architectural viewpoint. It is important to find an architecture, that is able to fulfill all requirements and is able to be evolved on for upcoming needs. In some areas where this is foreseeable, preparations can be implemented, whereas in other regions this is not possible. When working on a architecture, the SOLID **TODO: .** and YAGNI **TODO: .** principles can help to evaluate the stability and expandability **TODO: can it?**.

TODO: divide and conquer -> split issue in smaller issues, solve one by one

TODO: drafting, finding responsibilities, separate by responsibility

5.1 System Context

A System Context Diagram is suitable to find the boundaries of a system to implement and interactions with external components, therefore the first step in this analysis was to create Figure 5.1. The notation is a slight deviation from the UML standard so that the control and data flow can be displayed [2, p. 501].




Figure 5.1: System Context Diagram

This thesis will not focus on all parts mentioned in Figure 5.1, but for the implementation it is essential to have an as complete overview as possible.

5.2 Use Case Diagrams

A system context diagram is not suitable to elaborate all interactions in great detail, but use case diagrams can be used for that. They help to understand the customers needs [34] while the customer receives an impression on what will be reflected in the final product. Each but the very first of the following diagram focuses on one actor that is going to interact with the system. Figure 5.2 shows a high-level use case overview of all categories that are relevant to users of the system:




Figure 5.2: Use Case Diagram showing an high level overview of interactions

5.2.1 Managing Pipelines and Projects

One very important concern to the user of Winslow is to create an manage pipelines and to associate them to projects. For filtering and overview organization, the customer also expressed the need to attach tags to projects. This does not affect any execution strategy of Winslow, but regardless, it is still important to the end user.

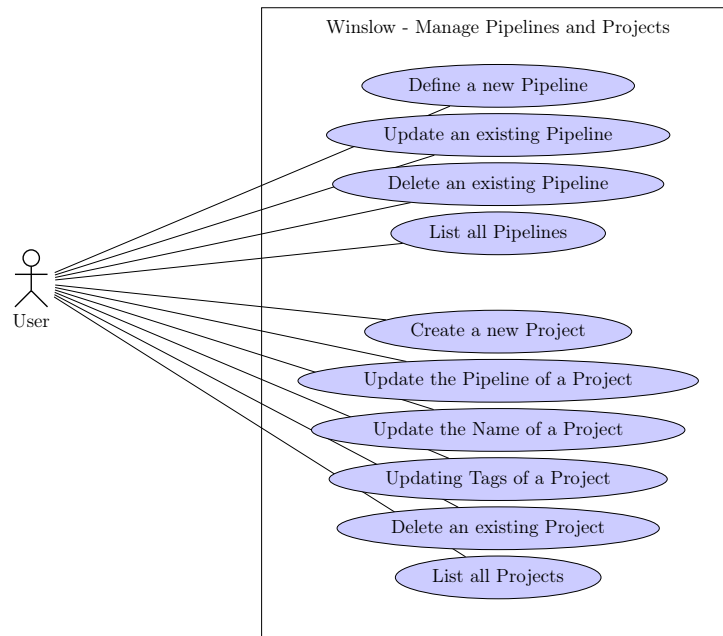


Figure 5.3: Use Case Diagram showing the general management interactions

5.2.2 Managing Resources and Workspaces

Uploading, listing and downloading files is one of the main concerns of Winslow to enable stage executions, so it is to no surprise that this is important for the user as well.




Figure 5.4: Use Case Diagram showing the general management interactions

5.2.3 Managing and Monitoring Executions

The main goal of the system is stage execution. The user expressed interest in advanced control, such as aborting a running stage or pausing the stage execution of a project pipeline, as well as being able to view logs live and inspecting intermediate output files .




Figure 5.5: Use Case Diagram showing the general management interactions

5.2.4 Monitoring Nodes

To monitor the usage of all nodes, an overview listing CPU usage, RAM usage and Network IO for each node was requested.




Figure 5.6: Use Case Diagram showing the general management interactions

5.2.5 System Administration

Furthermore to the interactions with a user, the system must provide further capabilities that are of concern to the administrator maintaining Winslow. What was raised here, is the appeal that the setup of new Winslow instances shall be easy as well as adding the new instance with an existing group of instances.

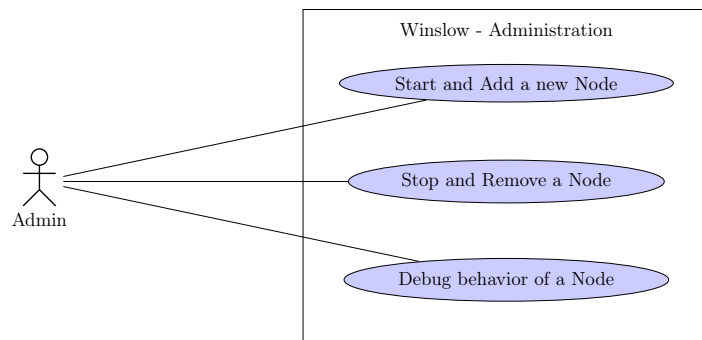


Figure 5.7: Use Case Diagram showing administrative interactions

5.3 Detailed Use Case analysis

The next step would be to write a detailed analysis for each use case, including information such as who is the initiator, which components are involved, what is the expected and what are alternative outcomes. But that would create further pages, with a lot of details for user interaction, which is not the focus for this thesis, which is the reason this is omitted here.

5.4 System Architecture

In Figure 5.8 a high level overview of the architecture is shown. It shows various services and regions with internal and external communication channels. The architecture follows a simplified¹ “Onion Architecture Pattern”[35], which is indicated by the layer numbers and colouring (inner layers are darker). For better readability, it is not arranged in circles.

TODO: mention user stories?

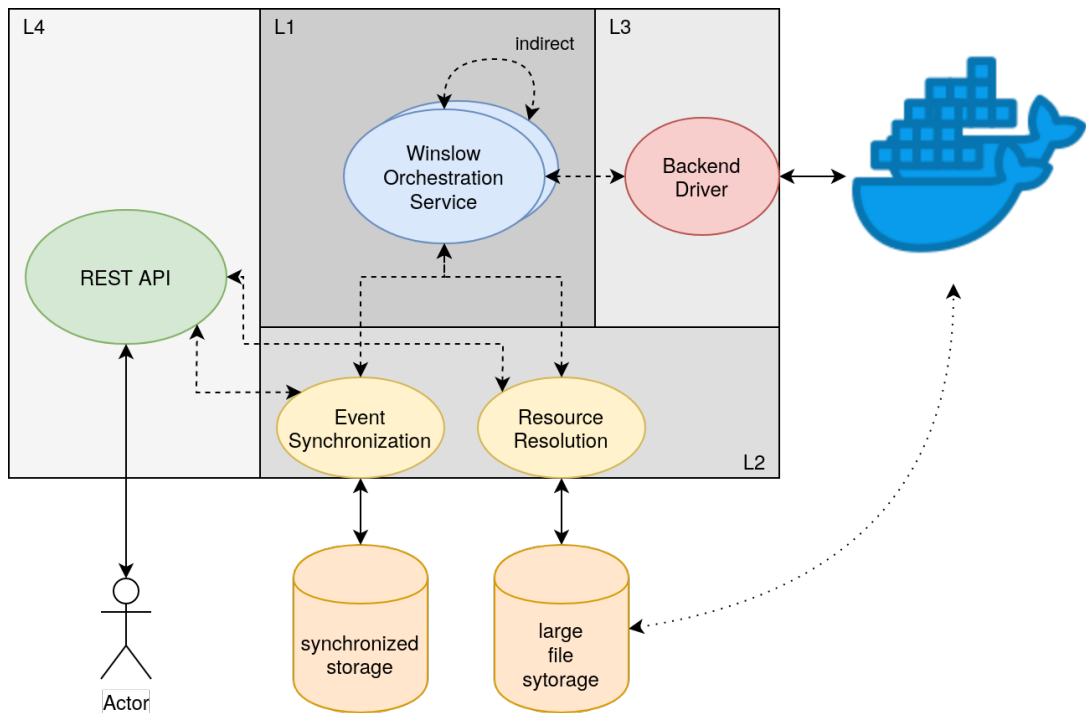


Figure 5.8: High level architecture overview of Winslow

The meaning of the layers and their services will be explained in the following sub-sections.

5.4.1 Layer 1: Orchestration Service

This layer is all about the fundamental business logic: when to schedule, start and fail which stage of what project as well as monitoring the hardware utilization.

¹fewer layers

This level is also responsible to communicate these actions with all other Winslow instances, so that there is no duplicate or missing stage execution or undetected node failure.

5.4.2 Layer 2: Events and Resources

The second layer is essential for the first layer to interact with and understand its environment. It provides two important services: a synchronized communication channel and a large file storage.

The communication channel requires a storage to persist messages that are relevant for a duration of time and it needs to support basic synchronization primitives to ensure consistency. Winslow instances that are started need to see events that were issued before their start and which are not completed yet to replicate the state of the system correctly.

For executing stages large files must be stored. The synchronized access is ensured by Winslow's event synchronization and thus the constraints to the storage is more relaxed than the one used to persist synchronization messages.

5.4.3 Layer 3: Backend Driver

This layer is responsible to interface with Docker. Once it is decided to run a stage this driver is instructed to start a certain image with environment variables and a prepared workspace. By separating this task to its own layer it will be easier to change from the Docker API to another implementation or away from the Docker platform at all if necessary. A reason for this could be the change of needs, the appearance of a platform that is fulfilling the needs better or the disappearance of the currently used platform. The latter sounds not that common at first, but the recent partial acquisition of Docker Inc by Mirantis[36] proves that even very popular third party software is not going to be around for all eternity. By moving the driver implementation into its own layer, but leaving the interface definition

in layer 1 allows the driver implementation to be swapped easily without the inner layer noticing. It then also complies to the “Dependency Inversion Principle”.

While the backend driver does not access any storage itself, the stages need to read and write to the large file storage. Docker provides means to mount local directories or remote NFS shares as volumes into the container natively. The access from within the executed stage can then be guarded to limit read and write access to only the required working directories.

5.4.4 Layer 4: Client Communication

The final layer in this architecture is the client communication layer. This service is not crucial for the actual execution and may be disabled on Winslow instances that shall not response to a user request themselves. A reason for this could be, that SSL certificate for a sub-domain points to a specific Winslow instances and unsecure access is not desired. The REST API shall provide resources that can be called from the statical served Angular Web-Application² or further tools to upload or download files, to control stage execution and to monitor usage and logs.

²**TODO:** see section 3.4

5.5 Time Schedule

The following figure shows schedule planned for this work:

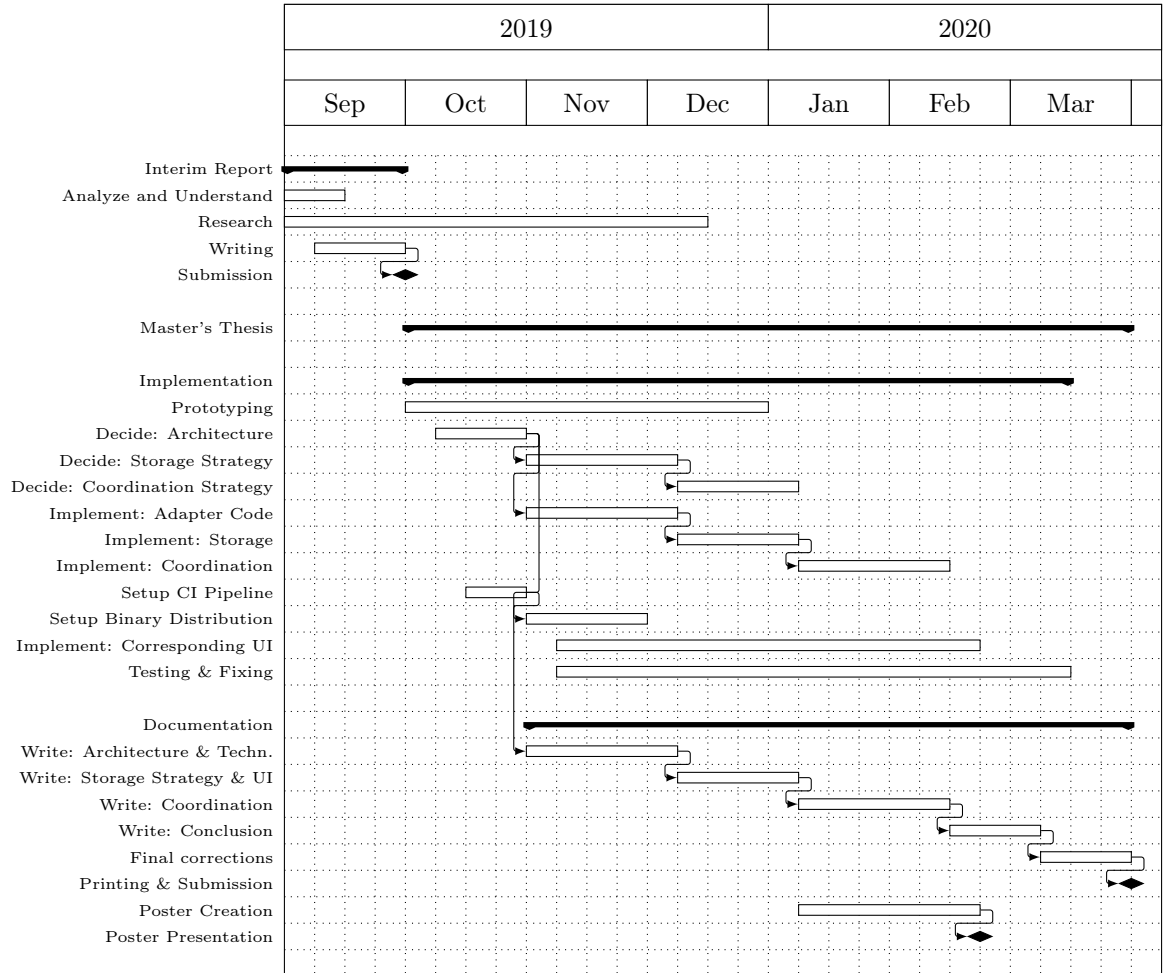


Figure 5.9: Time schedule

Chapter 6

System design

In this chapter the previously collected information are used to design the system in detail. The decision making process for that is documented here as well.

6.1 Environment

The environment of where the systems shall be deployed onto must be known for detailed design decisions, which will be elaborated here.

All Winslow instances are supposed to run inside a Docker container (as requested in section 2.2). The host will be a Ubuntu or Debian Linux Servers with some having multiple Nvidia graphics cards. There will be no graphical user interface available on these servers.

6.2 Storage Technology

Storing and accessing data is one of the main concerns for Winslow. In almost all use cases, there will be multiple Winslow instances that need to access the storage. Because of that, it must be accessible from remote and from multiple Winslow instances simultaneously. For that, the following technologies are reviewed: NFS or SMB/CIFS (subsection 3.3.1), GlusterFS (subsection 3.1.6), SeaweedFS (subsec-

tion 3.1.6), OpenIO (subsection 3.1.6), dCache (subsection 3.1.5) and HadoopFS (subsection 3.1.1).

The following criteria were used for comparison, with a scoring system with points from 0 to 3, in which 0 means the best solution and 3 the worst - the lower the overall score the better:

- **Installation Overhead:** How hard and extensive is the installation? Is there a good documentation? Is the package provided from within the main repository of Ubuntu, or alternatively, does it provide any other form of easy installation? Scoring 0 for installation through the main repository, 1 for an external repository, 2 for distribution specific archive, 3 for any other installation.
- **Dependencies:** Are further services required for operation? How hard are they to setup? Do they need any additional configuration? Scoring 0 for no further dependencies, 1 for dependencies but without the need of configuration, 2 for dependencies that require manual configuration, 3 for dependencies that need to be installed manually and require manual configuration.
- **Single Point of Failure:** Is the solution decentralized and is it failure resilient? Is it site- or rack-aware or provide replication mechanisms to compensate? Scoring 0 for completely decentralized, 1 for decentralized but with dependency on centralized backend, 2 for mostly centralized but with load balancing approaches, 3 for no single point of failure mitigations.
- **Docker integration:** How easy is it to integrate with Docker? Scoring 0 for native support, 1 for native but non-trivial support, 2 for support through additional exports facility, 3 for non-trivial solution.
- **Failure concerns:** Are there any noteworthy concerns? Was an early local test successful and reliable? Is it a known technology or “proven in

use”? Who is developing it and what are the support guarantees? Scoring 0 for commonly used and no concerns, 1 for minor concerns, 2 for major uncertainty, 3 for abandoned technology.

In Table 6.1 the results are displayed:

	NFS S.	SMB/CIFS S.	GlusterFS	SeaweedFs	dCache	HDFS
Inst.	0	1	0	2	3	1
Dep.	1	0	0	0	3	1
SPoF	2	3	0	2	0	0
Docker	0	2	2	3	2	3
Fail.	0	0	1	2	2	0
Score	3	6	3	9	10	5

Table 6.1: Comparison of storage technologies

The worst scoring tool in this comparison is dCache. The installation overhead, missing documentation and uncertainty on reliable operation is too high for this project (see subsection 3.1.5).

For similar reasons SeaweedFs is placed second worst. Local tests could not access SeaweedFs reliable when a node failed and there is no trivial support for Docker nor does it provide an NFS export. Using it from within Docker would require each container to be manipulated so that the first operation is mounting the storage with a custom binary and through a FUSE¹ mount. The tool also seems to be developed by a single person which introduces further uncertainty about reliability and support in the future.

HDFS and SMB/CIFS seem to be no terrible choice, but no especially good one either in this use case.

The two best scoring storage solutions are a simple NFS share and GlusterFS. While GlusterFS provides replication and decentralized access, a NFS share scores with its simplicity. The idea is to start with a plain NFS share for Winslow

¹Filesystem in Userspace

which can natively be utilized by Docker as volume mount and to revisit later whether the need for replication and decentralization persists. Because of the NFS interface export of GlusterFS, Winslow could then easily switch to GlusterFS by accessing the NFS export closest to its' instance.

6.3 Event Synchronization and Communication

kafka does what? persist events and provides them in an order(?), cant Winslow do this itself, so that no further dependency is required?

TODO: . <https://docs.microsoft.com/de-de/azure/architecture/microservices/design/intercommunication>

raw TCP (how to connect, centralized? star? tree?) centralized works towards -> centralized master node broadcast? microservice REST event bus? kafka messaging queue mqtt NFS provides filesystem, filesystems are kind of standardized

<https://kafka.apache.org/protocol> **TODO: .**

As discovered in subsection 5.4.2 there is a need for an event synchronization across Winslow instances. Without proper coordination, race conditions could cause stages to be started multiple times simultaneously, corrupt workspaces, configuration or project files. While starting too many stage executions are only wasting resources, data corruptions can lead to unrecoverable damage.

There are multiple ways to exchange data between systems that do not share the same process or machine. The most flexible implementation can be achieved by implementing a custom protocol on a raw TCP socket which allows to exchange blobs² between exactly two nodes. For a blob or message to reach all nodes, a connection to every other node, a centralized broker, another topology such as a tree structure or broadcast messages would be required. Because this in itself seemed to be a complex subject, third party alternatives were investigated.

Apache Kafka **TODO: CITE** is one such alternatives. It is open source, dis-

²**TODO: Binary Large Objects**

tributed and is focused on data and event streams. But it as every third party service, it introduces a dependency on the project. Not only in regards to interfacing and driver implementation, but also for maintenance and setup. Each Winslow Image would require to be shipped with a pre-configured Kafka service and the administrator must ensure that these services can reach each other - in addition to storage solution.

This was, where the idea appeared: Is it possible to not introduce any further dependencies to the system, but re-use the existing storage solution for synchronization and communication? To answer this, one must first investigate the smallest building block required for synchronized communication. Taking the idea from Kafka, an event stream, where the messages have are globally ordered: is it enough to provide synchronization over messages occurring in a synchronized sequence?

TODO: ACID in fundamentals? ACID, which stands for Atomicity, Consistency, Isolation and Durability, describes the needs a database must fulfill so that it does not suffer from side effects or data corruption when executing queries. A series of changes might not be allowed to be intercepted (atomicity), must see and result in consistent data, is not allowed to interfere with simultaneously executed queries and guarantees that once completed successfully, is in fact persistent and will not be lost (durability).

TODO: filesystem is close

TODO: DB principles: ACID

TODO: ...

TODO: requires common a way to exchange messages

TODO: all messages - called events - are published to all nodes

TODO: using file system as bus

TODO: clear and globally same order of events

TODO: events consist of an id, command, time, duration, subject and issuer

TODO: events might have a duration

TODO: unspoken requirement: all nodes share the same clock - or one with very little drift

TODO: usage as synchronization primitive

6.3.1 Messages

To coordinate the Winslow instances, messages will be exchanged through the common event bus. Because of the nature of this simple event bus, messages are delivered by broadcasting them to all instances.

TODO: issuer, command, subject, timestamp, duration

TODO: . Because this is a multi instance system which can suffer partial failure, all operations that are non-instantaneous require to have a timeout, so that a failure is detectable. Without a failure detection, the end of an operation might not be signaled, which could potentially block further operations for forever. The presence of a timeout will not be mentioned again when listing the messages in detail.

To account for transmission delay and clock offsets an additional time padding is granted. TODO: .

TODO: all messages have an id

TODO: to detailed here? move detailed description to system design?

- **LOCK**: Locking a named resource or operation. A lock is exclusive to the issuer and can only be granted if no other lock of the subject or parts of it exist. TODO: for example: a project needs to be locked to be modified by the issuer
- **EXTEND**: Extending an ongoing operation timeout. This can be used to signal, that a lock is required although its timeout is closing by, in which case the point in time the timeout is reached is pushed back. The issuer must be the same node as the one that issued the **LOCK** previously.

- **RELEASE**: An operation or resource is no longer required. The issuer must be the same node as the one that issued the **LOCK** previously.
- **KILL**: Non-gracefully stop an operation or destroy a lock. This signals that an operation on its own or associated to a lock shall be stopped immediately and that the lock shall be released. At the moment, this is only used whenever a user wants cancel a stage execution early. More details in [TODO: autoref](#).
- **ELECTION_START**: Signals that the need to execute a stage was detected. The project this signal refers to is included in the message.
- **ELECTION_PARTICIPATE**: Signals that the issuer node is capable of executing the stage. It also includes a scoring for the affinity and aversion of executing the stage on the issuer node (more details in section 6.8). This message also includes the project it refers to. This allows elections for different projects to take place concurrently.
- **ELECTION_STOP**: Signals that the election for a certain project has ended. The participator with the best affinity and aversion scoring is now allowed to start to execute the stage. More detail in [TODO: auref](#). This is signal is only allowed to be issued by the same node that started the election process.

[TODO: explain figure](#)




Figure 6.1: Sample event and lock series with lifetimes

6.4 Backend Driver

Docker REST API vs **TODO: Nomad** subsection 5.4.3

6.5 Execution Management

As noted in subsection 5.4.1, the central business logic is to deciding when and issuing stage executions. There are two general approaches in doing so, both will be discussed next.

6.5.1 Remote Execution

In this approach, the job is executed on a remote machine and not on the same node which has the responsibility in managing the execution. Continuous Integration (CI) platform Jenkins[5] as well as GitLab[4] do offer this approach. The so called slave node (Jenkins) or runner (GitLab) is accessed by the CI through a common interface (SSH in these cases) to start the job execution. Jenkins and GitLab even copy a custom binary onto the slave node, that is then managing the execution on the remote machine locally. This is due to the complexity in executing the job for a given configuration as well as being able to continue the execution on disconnects or maintenance reboots of the master machine.

In this scenario, the CI instance requires and stores login credentials for every remote machine to be able to login whenever needed. The system administrator therefore has to create a new user account on the remote machine, install required programs and environments (bot Jenkins and GitLab require an JRE³). In case of an security breach on the CI instance, the attacker is also able to login on all remote nodes and execute arbitrary programs as well **TODO: there is probably a great article about an issue like this!?**.

TODO: WRONG, on GitLab YOU need to intall/setup the runner

³Java Runtime Environment

GitLab follows another approach, the system administrator has to manually install the GitLab Runner⁴ on a slave node and is then able to add this runner to a Project. The security risk is somewhat similar to Jenkins: once an attacker has access to the CI instance, arbitrary code can be uploaded and executed on the remote machines. **TODO: but no ssh login until first binary uploaded**

TODO: winslow is more secure because of docker and therefore limited access to the system, no network..?

6.5.2 Local Execution

Executing a job locally means running a job on the same machine as the program that is monitoring the execution. This has the advantage of having all libraries, tools and resources already present **TODO: and is without the threat of connection loss to the remote machine**. But in the case of Jenkins and GitLab, this means, each CI instance is separate from the others. There is no integration between those instance in this configuration. What is missing Jenkins and GitLab here is the capability to decentralize their core task of managing projects, resources and executions.

TODO: scheduling

6.5.3 The approach for Winslow

TODO: .

potentially many executions -> easier when local

need watch hardware utilization closely

enables-ish “spin up a new instance and you are ready to go” slogan

6.6 Backend Nomad

TODO: .

⁴**TODO: cite**

6.7 Large File Storage organization

6.7.1 Stage storage

Thinking about the storage organisation for the pipeline and its stages, a few expectations and concerns arise. First of all, to redo a stage, one needs to be able to access the files that were the result of one, two or multiple stages before the current. Sometimes a stage wants to access intermediate data produces by multiple previous stages. Next, the input video footage needs to be accessed by multiple stages throughout the pipeline execution. Finally, some stage results are not intermediate but final results. **TODO: examples**

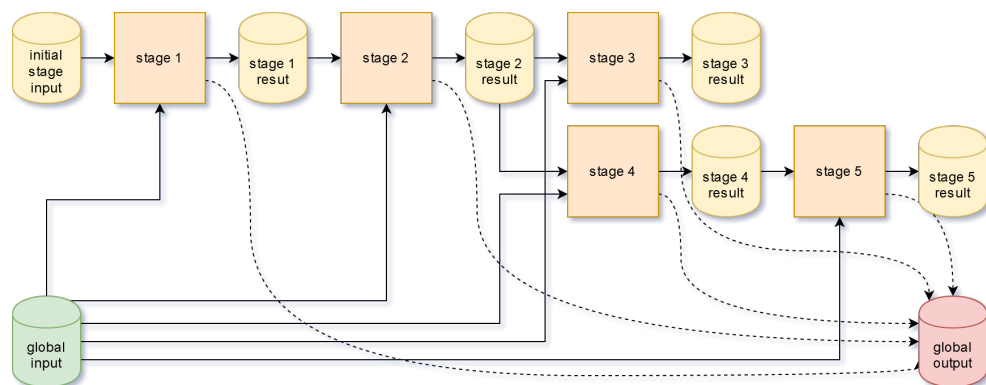
The first and second concern can be solved by providing a workspace directory for each stage, that is copied from the logically previous stage. Once the computation of a stage is completed, the workspace is considered immutable and only used to source new workspaces from. This works fine for small intermediate results, but it does not work very well for large files - like the video footage. Starting a new stage will take as long as the copy of multiple gigabytes on a spinning devices take (**TODO: sample GB and time**), require unnecessary storage due to multiple copies and provides no benefits in an archival and version control sense, because the video footage is not altered. So there needs to be another storage pool for input data, that is globally accessible and never changed: the global input storage pool. Providing one further storage pool for final results (global output pool), concern number three and four are also solved.

Because the very first stage has no workspace to source its files from, on creation of the pipeline a workspace directory for the “zeroths” stage is created. The user can then provide the very first stage with a predefined and non-empty workspace if necessary.

Because in the implementation of this project NFS was used, this also solves one further problem that was experienced in the prototyping phase: copying files on NFS (and many other network filesystems like Samba) is a client operation.

This means, the client reads the input file and writes to the output file. Because the common network used (gigabit in this case) is slower than spinning disk sequential read/writes (network: 112 Megabytes/s in theory, divided by two because read+write on the same connection: 56 Megabytes/s, spinning enterprise disk: about 170 Megabytes/s **TODO: test on johnny5**), this operation does not only take unnecessary long but also renders the network connection close to unusable for any other participant in the network that needs to use parts of the same physical route.

To ensure that the global input and the previous intermediate results are not altered, the Docker daemon is instructed to mount them as read-only filesystems. This makes it impossible⁵ for the stage to accidentally delete or alter the wrong files⁶.



In summary, the storage pool solution looks like this:

TODO: check paths

- `/input` is mounted read-only from
`nfs-server:/winslow/workspaces/<pipeline-id>/input`
- `/output` is mounted with write permissions from
`nfs-server:/winslow/workspaces/<pipeline-id>/output`

⁵When there is no bug in the implementation of mount or Docker

⁶Like in this scenario, where an unexpected space in a script caused people to loose their home directory: **TODO: steam rm -rf incident**

- `/workspace` is mounted with write permissions from

```
nfs-server:/winslow/workspaces/<pipeline-id>/<stage-id>
```

and was created by the copying the workspace directory of the logically previous stage

TODO: define logically previous stage -> previous stage on linear execution and ... when jumping around

initial input

stage execution does not need to depend on the result of the exact previous

6.8 Job Scheduling / Election System

TODO: .

TODO: affinity TODO: aversion

Utilizing Event Bus for timely limited elections and to ensure that there are no concurrent election for a single project.

TODO: what about concurrent elections on multiple projects

TODO: regarding locking: check for updatable, try to lock, check for updatable again, update, unlock

TODO: on unlock: check updatable

TODO: scheduling / job assignment/pull strategy

TODO: graph displaying over time lock and election

6.9 CPU RAM Network IO Monitoring

6.10 User Interface

TODO: REST request, read only: no lock; for changes: POST -> LOCK -> update -> RELEASE, on RELEASE Winslow gets notified and checks whether the new configuration leads to a new stage execution

6.11 File system

TODO: docker can auto mount nfs to container

6.11.1 Every node has a connection to every other node

6.11.2 Centralized broker

6.11.3 Tree hierarchy

6.11.4 outcome

6.12 Communication and Node Management

The system that shall be developed, is supposed to spread jobs onto execution nodes as available. There are two main approaches in node management and job assignment.

6.12.1 Centralized Management, Remote Execution

In the central management approach, there is always exactly one leader at any given moment in time. It is the responsibility of this leader to decide what to execute and where to execute it.

6.12.2 Decentralized Management, Local Execution

6.12.3 Combinations worth noting

stupid: - centralized management, local execution - decentralized management, remote execution

combination, decentralized + some remote slaves

6.12.4 Architecture

event based

common file system for communication because minimum requirements
stick to unix principles(list): simple, human readable intermediate format
voting/election by capabilities and 'will' of a node to run a stage

TODO: providing live hw utilization

TODO: the whole properties thingy

6.12.5 Failure handling

TODO: handling failed stages

TODO: handling failed nodes

6.13 Communication/Event architecture?

6.14 Targeting capabilities

6.14.1 General thoughts

6.15 Planned

6.16 Implementation details

6.17 Synchronization, Locking, publishing events

6.18 REST for UI

6.18.1 Atomicity of (Unix) Filesystems

6.18.2 Atomicity and behavior of NFS in particular

6.18.3 Using as lock backend

6.18.4 Using as election backend

6.19 Agile development

6.20 Continuous Deployment

Ironically, although mentioned as unusable in **TODO: .** for the use case Winslow wants to fulfil, it will use GitLab CI for continuous **TODO: delivery** of runnable Docker Images. Every code change in the Git repository will trigger a new build pipeline, that builds and tests Winslow, packages all dependencies **TODO: nomad** into a Docker image and pushes it to our in house private Docker

Registry **TODO: ..** The **TODO: execution nodes** will pull the latest image whenever they restart or are triggered to restart. **TODO: .**

Chapter 7

Outcome and Measurements

In this chapter the main concerns are listed. For each concern the current progress is described as well as further work that needs to be done.

TODO: compare to initial deliverable requirements

7.1 What did not work as expected

TODO: timeout for events might be reached when in between the bandwidth is used by a copy of a large file and the extend signal is because of that deferred, stages that are fine are then marked as failed

TODO: . old sections

7.2 Storage

One of the central concerns is the storage management. The program needs to make input files available on each execution node and collect the results once the computation is complete. There are a few main architectural strategies to approach this. Simplified, either at a centralized location which is accessed by all execution nodes, a copy of the input files to the execution nodes or decentralized and distributed between all execution nodes and replication. The advantages

and disadvantages can depend on the specific implementation and is therefore discussed in combination of such (see chapter 3).

Further testing is required to decide whether a more complex storage system is required, or the simplicity of a centralized solution outweighs the setup and maintenance overhead.

7.3 Coordination

Another important concern is the coordination of the nodes. A central coordinator with external server nodes, such as GitLab and Jenkins have, might not be sufficient for more complex and longer lasting pipelines. The probability that the master would need to be offline while there is a stage executed, is in the scenario of the desired workflow higher than for GitLab or Jenkins, because the stage is being execution for hours or days. Coupling stage execution plans on node availability ahead of time, as well as recovering from a sudden master failure implies additional implementation complexity. A decentralized coordination needs to be able to do this as well, but also allows the usage of the system while a node failed or is unreachable due to maintenance. With further prototyping and research a reasonable solution shall be found.

7.4 Binary distribution

In a time where containers are common and have proven to be usable, the installation of the binaries directly on the operating system they are executed on shall be avoided. There shall be no manual, nor automatic but custom file copies of the binaries or images from one server to the other. Experience shows, that without a proper management, this can easily become a mess, in which it is no longer clear, which files or images belong to which version. At the same time, making all binaries publicly available through the Docker Hub[37] is no option

either. Whether a self hosted Docker Registry[38] could be the solution to this will be determined in further testing.

7.5 User Interface

Providing a useful user interface might not be important to the functionality of the system itself but for the user experience. A bad user experience will cause a system not to be used. It became common practice for a rich user experience to be web based and interactive with JavaScript. For a potentially decentralized system, it is also advantageous to be able to access a disconnected node in the same manner as the remaining system, which further encourages a web based solution. Web based solutions such as React and Angular shall therefore be investigated for being used as user interface.

7.6 Requirements

TODO: more than first defined - as expected

TODO: lots more UI shortcuts found necessary while starting to use winslow

7.7 Architecture

TODO: loosely coupled backend driverFigure 5.8

7.8 High Level System Overview

TODO: intro words

TODO: explain what is going on in Figure 7.1




Figure 7.1: High level view on the system structure

7.9 Evaluation

The main focus of Winslow is to improve the efficiency in utilizing available hardware. Without it, new stages were seldom started at outside of the work time. To estimate whether Winslow is providing any improvements, the actions that are scheduled by Winslow outside of these hours are to be watched.

The following graphs are based on activity metrics, that were logged in the time period from the first usage in mid January 2019 until counter actions for COVID-19 in early March 2019 prevented active use due to the lock-down in Germany. For each day of the week and hour of the day an associated cell shows the activity colour encoded. Furthermore, red lines show the begin and end of a typical work day. Interestingly, the actual position of these red lines can also be inferred by the graphs themselves, due to activity spikes.




Figure 7.2: Actively running stages at a given day of the week and time of day

At first, Figure 7.2 and Figure 7.3 show the overall activity. In Figure 7.2 number of running stages per hour is shown and in Figure 7.3 the number of started stages per hour is shown. The increased activity between the read line - within the working hours - is visible. But especially in Figure 7.3 the activity

outside these times is visible as well.




Figure 7.3: Stages starting at a given day of the week and time of day

Figure 7.4 and Figure 7.5 show only stages that were triggered by Winslow, without a direct user input.




Figure 7.4: Actively running stages at a given day of the week and time of day which were started automatically

Figure 7.2 clearly shows that there is quite a bit of activity event without direct user input. This is especially obvious by looking at Figure 7.5: there are many stages started in the late evening and until midnight as well as a second wave at about one o'clock in the morning and activity on Saturdays. Without Winslow these stages would not have been triggered.

That the system was actively used to schedule work to be executed on the weekends. This can be seen in Figure 7.5 by highlighted areas late hours on Fridays, early and late hours on Saturdays as well as the continuous but decreasing overall activity from Saturday to Sunday in Figure 7.4. The gap of activity on Mondays at around noon in Figure 7.4 could indicate that the results were evaluated at here.




Figure 7.5: Stages automatically starting at a given day of the week and time of day

TODO: great staticists table

TODO: number of stages TODO: same graphs but where the stage was not manually started between 8h and (+9h) 17h TODO: table of number of stages executed, number of projects, manual trigger, auto trigger, average/min/max

time? TODO: percentage of stages triggered by user vs winslow




Figure 7.6: Stages automatically starting at a given day of the week and time of day

idle times that would be without automatically executing tasks

maybe with error scenario: being able to dynamically execute a stage of another project instead of stalling

TODO: hours spent in work time TODO: hours spent outside of work time

TODO: percentage win

Chapter 8

Further work

In the scope of this thesis a system was implemented to distribute the workload of a computer vision pipeline onto several compute nodes. Automatically stage scheduling improves the utilization of available hardware - especially at times where the staff is absent, like at nights or weekends.

Adding the support for WebSockets could help to make the Web-Application more responsive to serve side changes, like a stage state change. Pushing desktop notifications to the client or e-mail messages could reduce the response time of the staff for when a stage failed or a project needs attention otherwise. GlusterFs could be investigated for whether its (site) replication could reduce the usage of network bandwidth. **TODO: graphs** A system like this provides an endless stream of ideas to further improve the user experience or for extensions and capabilities to add.

Chapter 9

Conclusion

Overall the project was a success. Winslow helps to improve the efficient usage of available hardware resources by scheduling work in times, where no staff is watching. It helps in keeping order by separating projects and workspaces. The user interface provides an interaction method that is easy to understand for users, and provides upload and download mechanism without resorting back to console commands.

In the implementation phase an interesting and minimalistic way to synchronize events based on a commonly shared directory was discovered and implemented. Winslow is easy to setup and due to its few dependencies **TODO: no-mad, and kinda spring+angular** easy to maintain. Its architecture allows to be extended in further work. **TODO: .**

Bibliography

- [1] MEC-View Dr. Rüdiger W. Henn. *Mobile Edge Computing basierte Objekterkennung für hoch- und vollautomatisiertes Fahren*. URL: <http://mec-view.de/> (visited on 09/29/2019).
- [2] J. Goll. *Methoden des Software Engineering: Funktions-, daten-, objekt- und aspektororientiert entwickeln*. Springer Fachmedien Wiesbaden, 2012.
- [3] The Apache Software Foundation. *Apache Hadoop*. URL: <https://hadoop.apache.org/> (visited on 09/19/2019).
- [4] Inc. GitLab. *The first single application for the entire DevOps lifecycle*. URL: <https://about.gitlab.com/> (visited on 09/21/2019).
- [5] jenkins.io. *Jenkins. Build great things at any scale*. URL: <https://jenkins.io/> (visited on 09/19/2019).
- [6] Inc. GitLab. *GitLab CI/CD. Pipeline Configuration Reference*. URL: <https://docs.gitlab.com/ee/ci/yaml/> (visited on 09/19/2019).
- [7] jenkins.io. *Using a Jenkinsfile*. URL: <https://jenkins.io/doc/book/pipeline/jenkinsfile/> (visited on 09/19/2019).
- [8] Camunda Services GmbH. *Workflow and Decision Automation Platform*. URL: <https://camunda.com/> (visited on 09/19/2019).
- [9] Camunda Services GmbH. *Process Engine API*. URL: <https://docs.camunda.org/manual/7.6/user-guide/process-engine/process-engine-api> (visited on 09/19/2019).

- [10] Camunda Services GmbH. *Rest Api Reference*. URL: <https://docs.camunda.org/manual/7.8/reference/rest> (visited on 09/19/2019).
- [11] HashiCorp. *Nomad*. URL: <https://www.nomadproject.io/> (visited on 09/19/2019).
- [12] HashiCorp. *Nomad. Nomad vs. Kubernetes*. URL: <https://www.nomadproject.io/intro/vs/kubernetes.html> (visited on 09/19/2019).
- [13] Bc. Pavel Peroutka. *Web interface for the deployment and monitoring of Nomad jobs. Master's thesis*. URL: <https://dspace.cvut.cz/bitstream/handle/10467/80106/F8-DP-2019-Peroutka-Pavel-thesis.pdf> (visited on 09/22/2019).
- [14] Tigran Mkrtchyan Patrick Fuhrmann. *dCache. Scope of the project*. URL: <https://www.dcache.org> (visited on 09/19/2019).
- [15] Patrick Fuhrmann. *dCache, the Overview*. URL: <https://www.dcache.org/manuals/dcache-whitepaper-light.pdf> (visited on 09/19/2019).
- [16] Inc. Terracotta. *QuartzJob Scheduler*. URL: <http://www.quartz-scheduler.org/> (visited on 09/19/2019).
- [17] Data Revenue. *Distributed Python Machine Learning Pipelines*. URL: <https://www.datarevenue.com/en/blog/how-to-scale-your-machine-learning-pipeline> (visited on 09/19/2019).
- [18] Ask Solem. *Celery: Distributed Task Queue*. URL: <http://www.celeryproject.org> (visited on 09/19/2019).
- [19] IBM Knowledge Center. *IBM InfoSphere. DataStage*. URL: https://www.ibm.com/support/knowledgecenter/en/SSZJPZ_9.1.0/com.ibm.svg.im.iis.ds.design.doc/topics/c_ddesref_Server_Job_Stages_.html (visited on 09/19/2019).

- [20] Wikipedia contributors. *Qsub. Wikipedia, The Free Encyclopedia*. URL: <https://en.wikipedia.org/w/index.php?title=Qsub&oldid=745279355> (visited on 09/22/2019).
- [21] The University Of Iowa. *Basic Job Submission. HPC Documentation Home / Cluster Systems Documentation*. URL: <https://wiki.uiowa.edu/display/hpcdocs/Basic+Job+Submission> (visited on 09/22/2019).
- [22] Swiss National Supercomputing Centre. *High Throughput Scheduler*. URL: https://user.cscs.ch/tools/high_throughput/ (visited on 09/19/2019).
- [23] Colin Phipps. *zsync. Overview*. URL: <http://zsync.moria.org.uk/> (visited on 09/19/2019).
- [24] OpenIO. *High Performance Object Storage for Big Data and AI*. URL: <https://www.openio.io/> (visited on 09/22/2019).
- [25] Chris Lu. *Simple and highly scalable distributed file system*. URL: <https://github.com/chrislusf/seaweedfs> (visited on 09/22/2019).
- [26] Colin Phipps. *Alluxio. Data Orchestration for the Cloud*. URL: <https://www.alluxio.io> (visited on 09/19/2019).
- [27] Inc Red Hat. *Gluster. Free and open source scalable network filesystem*. URL: <https://www.gluster.org> (visited on 09/19/2019).
- [28] Inc. Docker. *Enterprise Container Platform*. URL: <https://www.docker.com/> (visited on 09/22/2019).
- [29] Inc. Docker. *Docker. Docker Logos and Photos*. URL: <https://www.docker.com/company/newsroom/media-resources> (visited on 01/16/2020).
- [30] Inc. Docker. *Docker Documentation. Swarm mode key concepts*. URL: <https://docs.docker.com/engine/swarm/key-concepts/> (visited on 03/17/2020).
- [31] The Kubernetes Authors. *Kubernetes. What is Kubernetes*. URL: <https://kubernetes.io/docs/concepts/overview/what-is-kubernetes/> (visited on 03/17/2020).

- [32] Inc. Docker. *Docker Documentation. Docker overview*. URL: <https://docs.docker.com/engine/docker-overview/> (visited on 01/16/2020).
- [33] IETF | Internet Engineering Task Force. *RFC 5661. Network File System (NFS) Version 4 Minor Version 1 Protocol*. URL: <https://tools.ietf.org/html/rfc5661%5C#section-1.7.2.2> (visited on 03/26/2020).
- [34] “Use Case Modeling”. In: *Use Case Driven Object Modeling with UML: Theory and Practice*. Berkeley, CA: Apress, 2007, pp. 49–82.
- [35] Jeffrey Palermo. *The Onion Architecture*. URL: <https://jeffreypalermo.com/2008/07/the-onion-architecture-part-1/> (visited on 03/26/2020).
- [36] Inc. Docker. *Press Release. Docker Restructures and Secures \$35 Million to Advance Developer Workflows for Modern Applications*. URL: <https://www.docker.com/press-release/docker-new-direction> (visited on 03/23/2020).
- [37] Inc. Docker. *Docker Hub. Build and Ship any Application Anywhere*. URL: <https://hub.docker.com/> (visited on 09/22/2019).
- [38] Inc. Docker. *Docker Documentation. Docker Registry*. URL: <https://docs.docker.com/registry/> (visited on 09/22/2019).