Brunel University London

College of Engineering, Design and Physical Sciences, Electronic and Computer Engineering

Distributed Computing Systems Engineering

Master Thesis

Conception and Realization of a Distributed and Automated Computer Vision Pipeline

---

Michael Watzko

1841795

March 31st, 2020

Declaration: I have read and I understand the MSc dissertation guidelines on plagiarism and cheating, and I certify that this submission fully complies with these guidelines.

College of Engineering, Design and Physical Sciences,
Electronic and Computer Engineering

Distributed Computing Systems Engineering

Master Thesis

Conception and Realization of a Distributed and Automated Computer Vision Pipeline

March 31st, 2020

A Dissertation submitted in partial fulfilment of the requirements for the degree
of Master of Science

Abstract

TODO: .polish

This thesis if about implementing a distributed system to process large work items with specialized hardware requirements in Docker container. The system furthermore assists the user by organizing the work items into projects, tracking execution progress and monitoring hardware utilization. All this is reported to the user in an Angular Web-Application which also allows the user to customize the execution pipeline, inspect intermediate results and to actively manage stage executions. Decentralized elections select the node which best fits the resources needs. The system and evaluation process is focused on easy setup and extendability.

Keywords: Events, Messaging, Synchronization, Software, Architecture, Distributed, Coordination, Docker, SpringBoot, REST, Angular, Typescript

Contents

Chapter 1

Introduction

Since the industrial revolution, humans strive for more automation in the industry as well as in the every day life. What was at first a cost saving measurement in factories, now also is a differentiation method for products. A new product must prove a higher level comfort to the customer than the previous generation as well as all the competitors. As such, the ambitions of the industry are focused on increasing the value of their products for the customer.

The automotive industry is one of the prime examples of this. Never was traveling from one place to another as comfortable as nowadays. Aspects like an elegant interior design, comfortable seats, air conditioning, entertainment systems and safety measurements need to be considered by car manufacturers to be competitive these days. The next luxury enhancement will be the autonomously driving vehicle. No longer shall the owner of a car steer it, but instead the car becomes his or hers personal chauffeur, driving the optimal route, the most comfortable way and being more reliable and safer than any human ever could.

The reason, autonomously driving cars are not common already, is their big complexity increase. Compared to already established technologies like parking assistants, entertainment systems or more efficient engine controllers, letting a computer reliably understand a certain traffic situation requires masses of input data and complex algorithms to process. As such, the problem itself becomes

massive and cannot be solved that easily. So the industry has no choice than to divide this into many small pieces and work out solutions step by step.

The MEC-View research project explorers one such step: whether and how to include external, steady mounted sensors in the decision finding process for partially autonomous vehicles in situations where onboard sensors are insufficient. To not disrupt traffic flow with non-human behavior, one needs to study and thereby watch human traffic. Automatically analyzing traffic from video footage requires a lot of computation power and can be further optimized by specialized hardware such as GPUs[1].

This thesis will conceptualize and realize a distributed and automated computer vision pipeline which is in this case used to analyzes traffic flow within video footage. Compared to an existing but highly manual workflow, the new system shall help to utilize the available hardware more efficiently by reducing idle times. Stage transitions and basic scheduling shall be automated to allow a user to plan and execute multiple projects ahead of time and in parallel.

---

[1]Graphics Processing Units

1.1 MEC-View

The MEC-View research project[1] - funded by the German Federal Ministry for Economic Affairs and Energy - aims to supplement the field of view of automated driving cars with road-side sensor data using 5G mobile communication. The sensor information is merged into an environment model on the so-called Mobile Edge Computing (MEC) server. This server is directly attached to the radio station to ensure low latency environment model updates.

The project is tested at an intersection in Ulm, Germany. Currently, there are 15 lidar and video sensors installed. Those sensors send their detections to the (MEC) server. A fusion-algorithm merges those detections into one environment model and sends it back to the (MEC) server and to the automated cars.

Additionally, general traffic flow is analyzed to learn about movement patterns. To do so, 4k video data is captured by an air drone from real world cross roads - not limited to the intersection in Ulm. On each frame of such a recording, cars are detected with a neuronal network. Detected cars are tracked throughout the video to compute the movement speed and position in time of each car. In an analysis of all vehicles, hot-spots of high and low traffic flow can be determined.

1.2 Focus

This thesis conceptualizes and implements a distributed and automated computer vision pipeline to increase the productivity in video analysis management and execution. It is not of concern for this thesis on how to retrieve the footage or what is further archived with the results of evaluated video footage. The focus is on utilizing available hardware resources, to manage multiple projects simultaneously and to reduce the idle time of relevant hardware by slow human reaction and availability times - like working hour constraints.

Chapter 2

Aims and Objectives

This chapter will discuss the program which shall be implemented. To do so, the problem to solve must be understood. To gather requirements and understand the technical hurdles to overcome, this chapter is split into two sections. First, a rough glance over the current workflow is given, which is followed by a more detailed description for the desired workflow.

2.1 Current Workflow

Currently, to analyze a video for the trajectories of recorded vehicles, the following steps are executed manually:

1. Upload the input video to a new directory on the GPU server

2. Execute a shell script with the video as input file and let it run (hours to days) until completed. The shell script invokes a Java Program - called TrackerApplication - with parameters on what to do with the input file and additional parameters.

3. The intermediate result with raw detection results is downloaded to the local machine and opened for inspection. If the detection error is too high,

the camera tracking has a drift or other disruptions are visible, the previous step is redone with adjusted parameters.

4. Upload the video and intermediate result to a generic computing server and run data cleanup and analysis. This is achieved with the same Java Program as in step 2, but with different stage environment parameters.

5. Download the results, recheck for consistency or obvious abnormalities. Depending on the result, redo step 2 or 4 with adjusted parameters again.

6. Depending on the assignment, steps 4 and 5 are repeated to incrementally accumulate all output data (such as statistics, diagrams and so on).

Because all those steps are done manually, the user needs to check for errors by oneself. Also, if a execution is finished or has failed early, there could be hours wasted until noticed, if the check intervals are too far apart, such as during nights or weekends.


Figure 2.1: Overview of workflow

5

2.2 Desired Workflow

The desired workflow shall be supported through an user interface that provides an overview of all active projects and their current state, such as running computation, awaiting user input, failed or succeeded.

To create a new project, a predefined pipeline definition shall be selected as well as a name chosen. Because only a handful of different pipeline definitions are expected, the creation of such does not need to happen through the user interface. Instead, it is acceptable to have to manually edit a configuration file in such rare circumstances.

Once a project is created, the user wants to select the path to the input video. This file has to be been uploaded to a global resource pool at this point. The upload and download of files shall therefore also be possible through the user interface. Because a video is usually recorded in 4k (3840 x 2160 pixels), encoded with H.264 and up to 20 minutes long, the upload must be capable of handling files which are tens of gigabytes large.

Once a pipeline is started, it shall execute the stages on the most fitting server node until finished, failed or a user input is required. Throughout, the logs of the current and previous stage shall be accessible as well as uploading or downloading files from the current or previous stages workspace. In addition to the pipeline pausing itself for user input, the user shall be able to request the pipeline to pause after the current stage at any moment. When resuming the pipeline, the user might want to overwrite the starting point to, for example, redo the latest stage.

Mechanisms for fault tolerance shall detect unexpected program errors or failures of server nodes. Server nodes shall be easily installed and added to the existing network of server nodes. Each server node might provide additional hardware (such as GPUs), which shall be detected and provided.

For the ease of installation and binary distribution, Docker Images shall be used for running the Java Program for analyzing the videos as well the to be implemented management software.

2.3   Deliverable Requirements

From the desired workflow, the following requirements can be extracted:

- User interface for interaction between the system and the user

- Storage management for global resource files as well as stage based workspaces

- Pipeline definition through configuration files

- Handling of multiple projects with independent progress and environment

- Reflecting the correct project state (running, failed, succeeded, paused)

- Log accumulation and archiving

- Accepting user input to update environment variables, resuming and pausing projects as well as uploading and downloading files into or from the global resource pool or a stages workspace.

- Assigning starting stages to the most fitting server node

- Detecting program errors (in a stage execution)

- Cope with server node failures

- Providing a Docker Image for the implemented program, preferably in an automated fashion.

Furthermore, a given constraint is that the user interface is to be implemented as Angular Web-Application using a REST API (for more details see section 3.4).

While fulfilling these requirements, the overall goal of the system is to increase productivity and hardware utilization, especially in times where the staff is absence. The need to constantly monitor and interfere with the system shall be decreased by providing automated mechanisms where possible and appropriate.

2.3.1 Non-Requirements

To know the requirements and expectations of a system is essential, but knowing what is not expected by the system is at least as valuable. It prevents wasting resources, effort and architectural specialization that will never be required or in the worst case, make further development harder by restricting available choices in the future.

The system to implement shall not strive to implement real time scheduling or low latency scheduling. The expected work items are big chunks that require hours to compute. Whether the assignment of the work item takes sub seconds or several seconds is nearly unnoticeable in the overall compute time.

Chapter 3

Literature Survey

This chapter is about accumulating information. Programs that are solving similar problems, as described in the desired workflow, or dealing with a subset of the problem are looked into. Fundamental knowledge for understanding this thesis is also acquired here.

3.1   Similar solutions

This sections focuses on programs trying to provide somewhat similar workflows. The reason for this is to use well established or suitable programs as middle-ware to reduce implementation overhead, or, where this is not possible, one might be able to gather ideas and learn about proven strategies to use or pitfalls to avoid while implementing custom solutions.

3.1.1   Hadoop MapReduce

For big data transformation, Hadoop MapReduce[2] is well known. With MapReduce, the input data is split into blocks of data and distributed onto Mappers. Mappers then apply the business logic and output intermediate results in form of key/value pairs. After shuffling, the Reduce stage will combine values from key/-value pairs with the same key. In the final output, each key is unique accompanied

with a value.

This strategy has proven to be very powerful to process large amount input data because Mappers and Reducers can work independently on their data-sets and therefore scale very well when adding further instances.

If the implementation were to be based on Hadoop MapReduce to achieve the desired workflow, it could be done like the following:

- Each video is split into many frames and each frame is applied to a Mapper

- A Mapper tries to detect all vehicles on a frame and outputs their position, orientation, size and so on

- The Reducer then tries to link the detections of a vehicle through multiple frames

- The final result would be a set of detections and therefore all positions for each vehicle in the video

But at the moment, this approach seems to be unfitting due to at least the following reasons:

1. It is not always trivial to reasonable link the detections of a vehicle. For example, a vehicle can be hidden behind a tree for a few frames until visible again. In addition, MapReduce requires the combination to be performed per common key. Until one is trying to link the detections of multiple frames, there is no common identifier that could easily be used as key. The position of a moving vehicle cannot be used as key, neither can the color or size, because of the noise of the camera, deviation in detection output and perspective distortions. The current implementation of the TrackerApplication is archiving this by finding similarities between detections, but for the Mapper it would be required to express this as a deterministic key.

2. MapReduce is great in combining many machines to solve a big computational problem. But at the moment, this is neither a desired nor given

condition. At the moment, there are a handful of very powerful workstations with specialized hardware. Therefore it is perfectly acceptable and sometimes required, when each workstation works through a complete video at a given time instead.

3.1.2 Build Pipelines

Build pipelines such as GitLab[3] and Jenkins[4] can also distribute the execution of stages onto other server nodes. In a common use-case, such build pipelines are used to build binaries out of source code, after a new commit into a SCM[1] repository was made. At IT-Designers GmbH GitLab as well as Jenkins are commonly used for scenarios exactly like this. A pipeline definition in GitLab CI/CD [5] or in a Jenkinsfile [6] describe stages and commands to execute. Each stage can be hosted on another node and be executed sequential or in parallel to each other.

Although this seems to be quite fitting for the desired workflow, there are two issues. First of all, such a pipeline does not involve any user input besides an optional manual start invocation. The result is then determined based on the state of the input repository. Second, such a pipeline is designed to determine the output (usually by compiling) whereas each run is independent from the previous and a repeated run shall provide the same result as the previous did. Usually, a new run is only caused by a change of the input data. However, the desired workflow differs in this aspects. A redo of a stage can depend on the result of the previous stage, for example, if the results are poor or the the stage failed. Instead of having multiple complete pipeline runs per project, the desired workflow uses a pipeline definition as base for which the order can be changed. Also, intermediate results need to influence further stages, even if repeated.

---

[1]Source Code Management

3.1.3   Camunda

Camunda[7] calls itself a "Rich Business Process Management tool" and allows the user to easily create new pipelines by combining existing tasks with many triggers and custom transitions. Camunda is focused upon visualizing the flow and tracking the data through a pipeline. The Camundas Process Engine[8] also allows user intervention between tasks.

One of the main supporting reason for it Camunda is the out of the box rich graphical user interface for process definition and interaction. Through its API[9], Camunda also allows custom external workers to execute a task. But it misses the capability to control which task shall be processed on which worker node which is required by the desired workflow. It does also not provide any concept on how to allocate and distribute resources. The user interface - while being rich overall - is quite rudimentary when it is about configuring tasks and would therefore require custom plugins to be developed for more advanced user interactions.

Camunda is also not designed to reorder stages or insert user interactions at seemingly random fashion. The user itself is considered more as a worker that gets some request, "executes" this externally and finally marks the request as accepted or declined. Mapping this to the desired workflow does not feel intuitive. Finally, there is also no overview of task executors, no centralized log accumulation and no file up- or download for global project resources.

3.1.4   Nomad

Nomad[10] by HashiCorp is a tool to deploy, manage and monitor containers, whereas each job is executed in its own container. It provides a rich REST API and can consider hardware constraints on job submissions. Compared to Kubernetes[11], which is similar but more focused on scaling containers to an externally applied load, it is very lightweight. It is also available in many Linux software repositories - such as for Debian - which makes the installation very easy.

Because there were no grave disadvantages found (depending on a third party

library can always be considered be a disadvantage for flexibility, error-pronous and limit functionality) Nomad is being considered as a middle-ware to manage and deploy stages. Others[12] seem to be using Nomad to manage and deploy containers for similar reasons. Nonetheless, further testing and prototyping will be required for a final decision.

3.1.5   dCache

"The goal of this project is to provide a system for storing and retrieving huge amounts of data, distributed among a large number of heterogenous server nodes, under a single virtual filesystem tree with a variety of standard access methods"[13]. dCache seems to be able to solve the storage access and distribution concern for the stages and sever nodes. When using dCache, one could store the global resources distributed between the server nodes. Built-in replication would prevent access loss on a node or network failure and an export through NFS[2] allows easy access for Linux based systems[14].

But the installation is complex and requires many services to be setup correctly, such as postgresql and many internal services such as zookeeper, admin, poolmanager, spacemanager, pnfsmanager, cleaner, gplazma, pinmanager, topo, info and nfs. The documentation is also rather outdated and incomplete which meant, early tests with a prototype setup took days to setup and behaved rather unstable (probably due to a wrong configuration). It is to be seen, whether such an complex and heavy system is actually required or if there are feasible alternatives.

3.1.6   Further mentions

The following list acknowledges programs that behave similar to the previously mentioned strategies or can be used as building blocks for custom implementations. Programs that are listed here, were looked into, but not all in great depth

---

[2]Network File System

because miss-fits or missing functionality were detected early on. The list is in no specific order:

- **Quartz**[15] is a Java based program to schedule jobs. Instead of doing so by using input, Quartz executes programs through a timetable and in certain intervals.

- **Luigi**[16] also executes pipelines with stages and is written in python. The advertised advantage is to define the pipeline directly in python code. But, this is at the same time the only way to define pipelines which contradictions with the existing Java TrackerApplication implementation.

- **Celery**[17] is focused on task execution through message passing and is written in Python. Intermediate results are expected to be transmitted through messages. Because there is no storage strategy and python adapter-code would have been required, Celery was dismissed.

- **IBM InfoSphere**[18] provides similar to Camunda a rich graphical user interface but for data transfer. Dismissed due to commercial nature.

- **qsub**[19][20] is a CLI[3] used in HPC to submit jobs onto a cluster or grid. Dismissed due to an expected high setup overhead, non-required multi-user nature and the fact, that it only provides a way to submit jobs.

- **CSCS**[21] High Throughput Scheduler (GREASY). Dismissed for similar reasons as qsub, although it is more light weight and hardware agnostic (it can consider CUDA/GPU requirements).

- **zsync**[22], similar to rsync, is a file transfer program. Zsync allows to only transfer new parts when a file that shall be copied already exists in an older version on the target. This tool might be useful when implementing a custom resource distribution strategy is required.

---

[3]Command Line Interface

- **OpenIO**[23] provides a distributed file system, is already provided as Docker image and provides a simple to use CLI. Because the NFS export is only available through a paid subscription plan, it was dismissed from further investigation.

- **SeaweedFS**[24] provides a scalable and distributed file system. The most interesting aspects are that it is rack-aware as well as natively supports external storage such as Amazon S3. When adding server nodes from the cloud this could allow all nodes to access the same file system while using rack-aware replication to reduce bandwidth usage and latency. A local test also proved that it is easy to setup, but because it cannot hot-swap nodes and was not able to recover when the seeweed master node became unreachable it was dismissed.

- **Alluxio**[25] provides a distributed file system but was dismissed because it itself requires a centralized file system for the master and its fallback instances

- **GlusterFS**[26] is another tool to provide a distributed file system with replication. It was bought by IBM but is nonetheless available through the software repository of many Linux distributions such as Debian. A local test showed that the setup is very easy and no adjustments of configuration files are required. However, the replication mechanism requires that an integer multiple of nodes of the replica value are assigned to the file system. This makes GlusterFS hard to use in a scenario, where adding and removing nodes are expected to happen frequently.

3.2   Docker

As describe in section 2.2, for easy deployment and setup Docker[27] containers shall be used. This allows easier isolation of the executed stages and other host

programs, as will be explain further on.

Docker is the name of a software that combines various isolation technologies (subsection 3.2.1) to provide and execute third party software in virtual environments. Docker aims are to increase security and to simplify installation as well as maintenance of applications. Docker Swarm[29] and Kubernetes[30] use these fundamental features to scale applications, services and microservices in the cloud. Software is retrieved from a public registry, called Docker Hub[31] or from local or private registries.


Figure 3.1: Official Docker "Moby" Logo[28]

3.2.1 Technology

Docker uses so called images to package and transport binaries with all their required libraries and configuration files as a read-only archive to the destination system. Instead of spawning a new process for a binary in the host environment, the image is used to create a new virtual environment - the so called container. Because the image includes all required libraries, it is possible to completely isolate the virtual environment which makes the host system invisible to any process inside the container. Changes to files within containers are stored separately as differentials[4], which allows an image to be used by multiple containers at once. Privileges, resource limitations and storage configurations can also be specified when creating a new container. Processes inside containers are unable to see other processes or files that are not part of or assigned to the their container[5].

In contrary to a hypervisor, docker is archiving this without hardware and operating system virtualization. Instead of executing the container inside an additional virtualized operating system, they share the kernel of the host system.

---

[4]for example by using OverlayFS[32]

[5]This is the default behaviour. It is possible to manually lift or modify many boundaries Docker enforces for containers on default.

For that, the host needs to support additional isolation mechanisms At the time of writing this, only the Linux Kernel is capable to separate processes, network interfaces, interprocess communications, filesystem mounts and the time-sharing systems by namespaces. By configuring these namespaces, Docker is capable to isolate containers into virtual environments. Furthermore, control groups can be used to limit and constraint access to hardware resources. [33]

These approaches allow containers to run with very little additional overhead in comparison to running the application directly on the host. Containerization increases security by limiting what an application sees and is able to interact with, decreases maintenance overhead because of no additional operation systems to maintain and allows to run multiple instances of the same application besides each other with independent configurations and environments but with the same base image.

3.3 Data Storage

When distributing workload onto different machines, accessing input and output files becomes another concern which is not present when all computation and storage resides on the same physical machine. In theory, one could use portable mediums such as USB-Sticks, CD-ROMs or external HDDs, but in reality, this becomes very tedious really fast. More advanced users might be able to take advantage of a common network connection between the computation to copy files and directories to the required places. This strategy - still being tedious - surfaces another issue: dealing with multiple copies requires careful version management and additional storage space. One certainly would not want to continue computation on outdated data or have too many copies of the same files.

The solution to that can be a network file system. To programs this solution seems to be just another local directory hierarchy, but in reality, the files might

be located on another or multiple other machines.

3.3.1 Remote File System

Remote file systems provide access to data that does not need to be stored on the local machine and can often be shared with other clients as well. Some file systems are organized centralized, where at a single address and physical location all data is stored and retrieved from. This can have advantages in setup time, maintenance and communication complexity, because of their simplicity.

The Network File Systems (NFS) for example, which is primarily used in the Linux and Unix world, is centralized, the clients are connecting to a remote server to read and write files from and to. Since NFS 4.1 pNFS[6][34, p. 14, section 1.7.2.2] allows the server to spread load across multiple servers and volumes, but because all metadata is still handled by one server, this approach still has a single point of failure. The Common Internet File System (SMB/CIFS) is similar to NFS in regards of providing access to a remote file system locally, and it is the approach primarily used by Windows computers. Similar to NFS, it has also a centralized approach.

More advanced file systems provide additional functionality like replication, which stores the data more than once to recover from data carrier failures, site awareness, which replicates data to geographically distant locations[7] and decentralized communication endpoints that provide access to the data from more than a single point of failure.

File systems like GlusterFS (subsection 3.1.6), Alluxio (subsection 3.1.6), SeaweedFS (subsection 3.1.6), OpenIO (subsection 3.1.6), dCache (subsection 3.1.5) and HadoopFS (subsection 3.1.1) try to provide these functionality.

---

[6]parallel NFS

[7]to not be affected by a burning data centre for example or to decrease access time for clients from the distant location

3.4 Angular Web-Application and REST API

The user interface is supposed to be implemented by an Angular Web-Application. Angular[35] itself is a framework for mobile and desktop Web-Applications which uses templating to generate the website completely on the client. It uses Type-Script[36], which is, in essence, an enriched subset of JavaScript with type annotations, that compiles back to JavaScript. This allows any recent web-browser to execute TypeScript "code", without the need to explicitly supporting it. The additional type annotation in TypeScript supports developers by giving helpful error messages, instead of runtime errors.

This thesis uses REST REST (Representation State Transfer) to enable the Angular application to communicate with the web-server. REST is, by now, commonly used to transfer state on top of HTTP in the internet. It utilizes the HTTP methods[37] to request resources or to push updates using URLs for localization.

TODO: cleanup

Chapter 4

Introducing Winslow

The program that shall implement the requested features listed by **??** is being called Winslow. This name refers to Frederick Winslow Taylor who was an American mechanical engineer and one of the first management consultants in the 19th century[.] Both strive to make people's work more efficient.


Figure 4.1: Winslow Logo

4.1   Common Terminology

This chapter explains the meaning of words and terminology used later on. It is crucial that every reader has the same understanding for the words used so that there are no wrong assumptions, expectations or surprises.

The root of a work item is called project. A project will usually refer to one video footage that shall be processed. Each project has its exclusive directory for input, output and intermediate files, called workspace. Furthermore a project as an author and possibly participants, that can help in steering and monitoring the processing. To do so, there is always a pipeline assigned to a project.

A pipeline consists of at least one but usually multiple stages that can at least be processed linearly - one after the other. Furthermore, a pipeline can define environment variables that valid for all stages within the stage.

A stage is the smallest work unit that can be executed. A docker imagesection 3.2 is specified as execution environment as well as further stage specific environment variables, command line arguments and hardware requirements (such as CPU cores, GPUs and minimum available RAM). This enables a stage to use a common image as well as to specify very precisely how to process the data.

In addition, stages and pipelines can be distinguished in a active, running or completed stage or pipeline and a definition. A stage definition specifies the above mentioned presets - allowing the user to adjust details just before submission - to create multiple stages from, while a running or complete stage has all information of what is or was executed and allows no further alteration. For pipeline definitions this is similar, a pipeline definition can be used to specify a common order of execution. Once assigned to a project, the project can freely adjust and change its instance of the pipeline.

A running process that is executing the binary of Winslow is called Winslow instance.

A computer with a Winslow instance that is accepting or executing stages is called execution node. The focus relies here on the computer providing compute resources to Winslow.

TODO: cleanup

Chapter 5

System Analysis

This chapter is about analysing the to be implemented system. It is important to find an architecture that is able to fulfil all requirements and that is able to be evolved for upcoming needs. But, as the YAGNI[38, p. 36] principle describes, this does not mean one should plan for speculative events. In some areas where changes or extensions are foreseeable, preparations can be considered, whereas if there is no such indication, one should only consider the actual and current need. In addition, keeping the SOLID[39, p. 86] patterns in mind can help to improve correctness, simplicity, stability, changeability and testability[38, p. 162].

The analysis and design follows the "Top-Down" approach[40, p. 171] in which the system is first outlined and then step by step partitioned into smaller and more detailed pieces.

5.1 System Context

A System Context Diagram is suitable to find the boundaries of a system to implement and interactions with external components, therefore the first step in this analysis was to create Figure 5.1. The notation is a slight deviation from the UML standard so that the control and data flow can be displayed [40, p. 501].
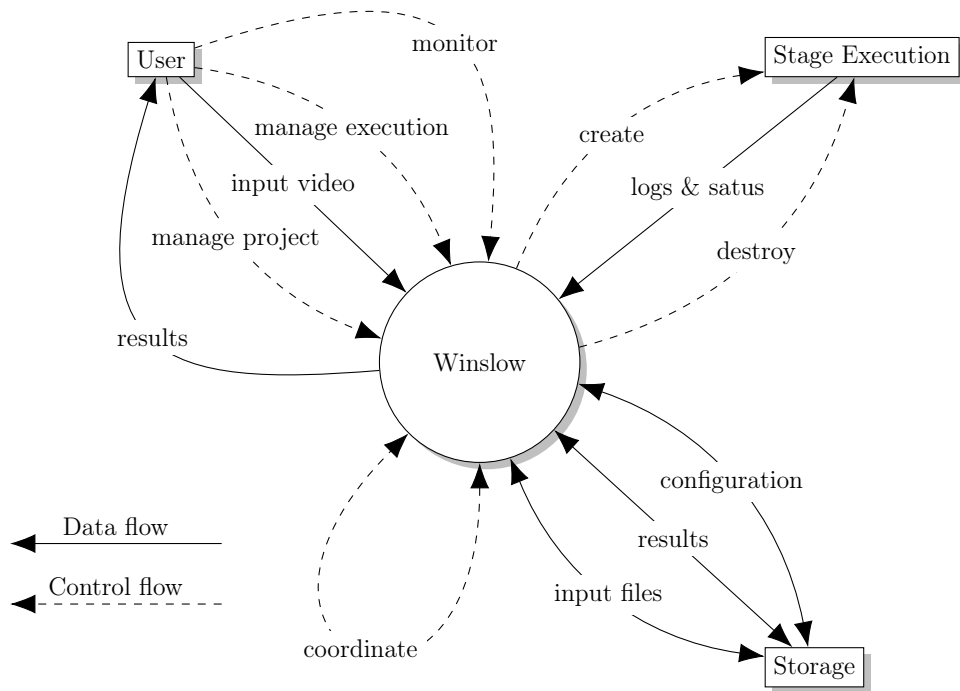

Figure 5.1: System Context Diagram

This thesis will not focus on all parts mentioned in Figure 5.1, but for the implementation it is essential to have an as complete overview as possible.

5.2 Use Case Diagrams

A system context diagram is not suitable to elaborate all interactions in great
detail, but use case diagrams can be used for that. They help to understand the
customers needs [41] while the customer receives an impression on what will be
reflected in the final product. Each but the very first of the following diagram
focuses on one actor that is going to interact with the system. Figure 5.2 shows
a high-level use case overview of all categories that are relevant to users of the
system:
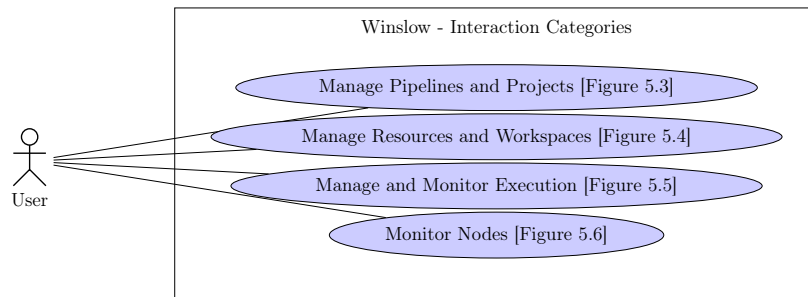
Figure 5.2: Use Case Diagram showing an high level overview of interactions

5.2.1 Managing Pipelines and Projects

One very important concern to the user of Winslow is to create an manage
pipelines and to associate them to projects. For filtering and overview organiza-
tion, the customer also expressed the need to attach tags to projects. This does
not affect any execution strategy of Winslow, but regardless, it is still important
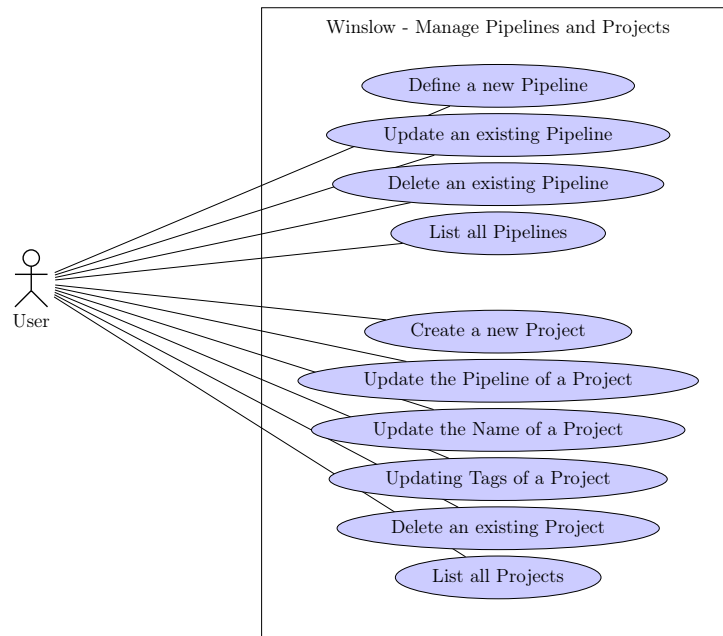to the end user.

Figure 5.3: Use Case Diagram showing the general management interactions

5.2.2 Managing Resources and Workspaces

Uploading, listing and downloading files is one of the main concerns of Winslow to enable stage executions, so it is to no surprise that this is important for the user as well.
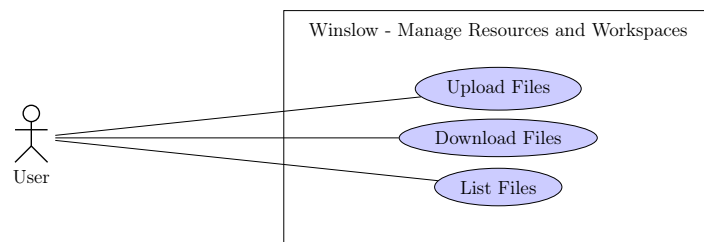
Figure 5.4: Use Case Diagram showing the general management interactions

5.2.3 Managing and Monitoring Executions

The main goal of the system is stage execution. The user expressed interest in advanced control, such as aborting a running stage or pausing the stage execution of a project pipeline, as well as being able to view logs live and inspecting intermediate output files .

Figure 5.5: Use Case Diagram showing the general management interactions

5.2.4 Monitoring Nodes

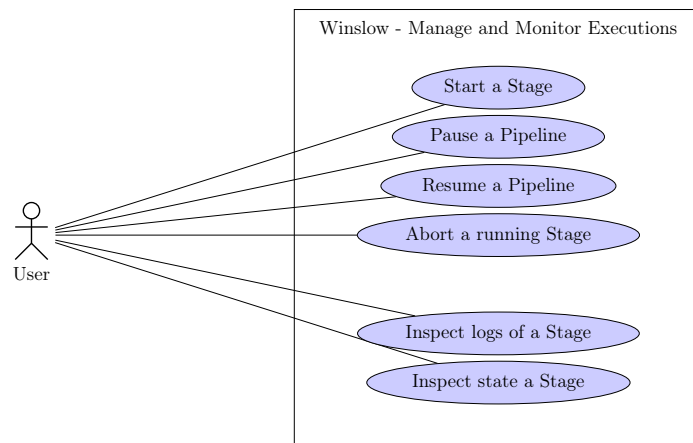To monitor the usage of all nodes, an overview listing CPU usage, RAM usage and Network IO for each node was requested.

Figure 5.6: Use Case Diagram showing the general management interactions

5.2.5 System Administration

Furthermore to the interactions with a user, the system must provide further capabilities that are of concern to the administrator maintaining Winslow. What was raised here, is the appeal that the setup of new Winslow instances shall be easy as well as adding the new instance with an existing group of instances.
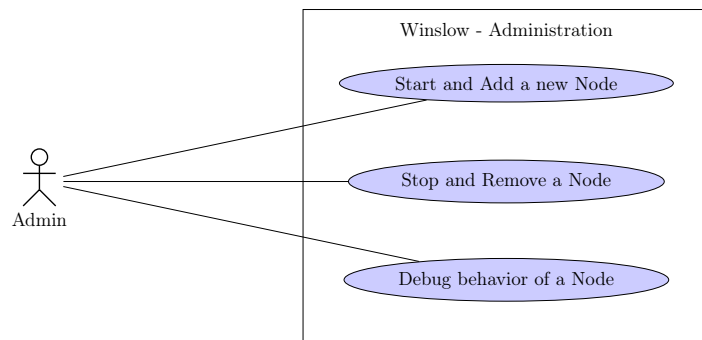
Figure 5.7: Use Case Diagram showing administrative interactions

5.3 Detailed Use Case analysis

The next step would be to write a detailed analysis for each use case, including information such as who is the initiator, which components are involved, what is the expected and what are alternative outcomes. But that would create further pages, with a lot of details for user interaction, which is not the focus for this thesis, which is the reason this is omitted here.

5.4 System Architecture

In Figure 5.8 a high level overview of the architecture is shown. It shows various services and regions with internal and external communication channels. The architecture follows a simplified[1] "Onion Architecture Pattern"[42], which is indicated by the layer numbers and colouring (inner layers are darker). For better readability, it is not arranged in circles.

Figure 5.8: High level architecture overview of Winslow

The meaning of the layers and their services will be explained in the following sub-sections.

5.4.1 Layer 1: Orchestration Service

This layer is all about the fundamental business logic: when to schedule, start and fail which stage of what project a well as monitoring the hardware utilization.
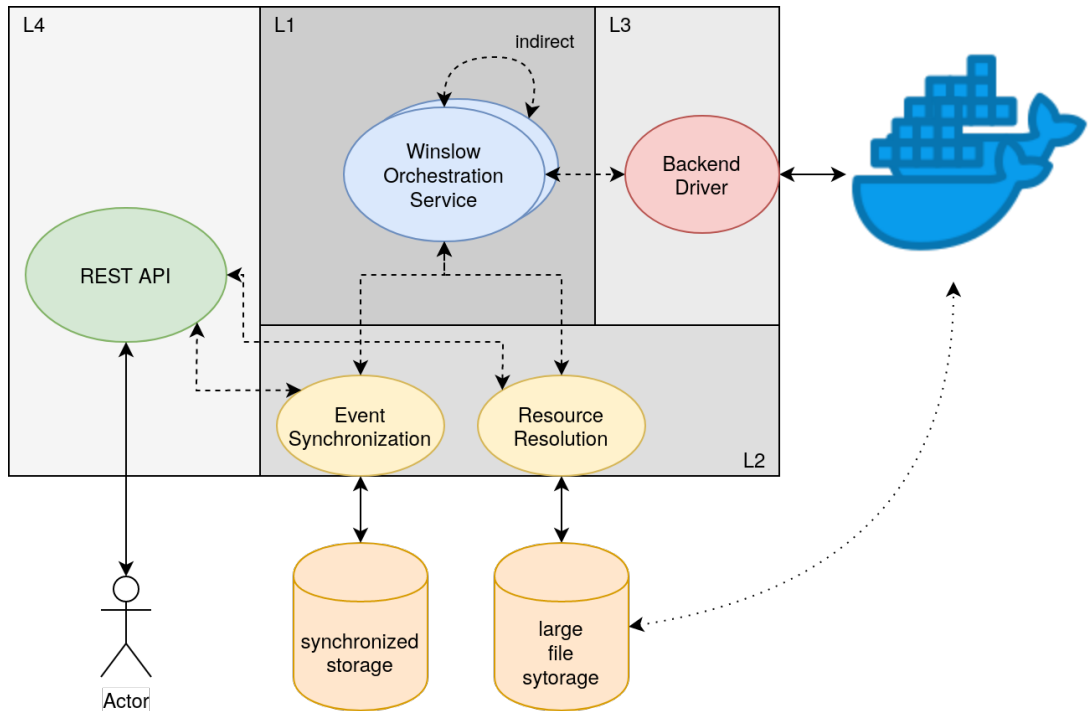
---

[1]fewer, collapsed layers

This level is also responsible to communicate these actions with all other Winslow instances, so that there is no duplicate or missing stage execution or undetected node failure.

5.4.2   Layer 2: Events and Resources

The second layer is essential for the first layer to interact with and understand its environment. It provides two important services: a synchronized communication channel and a large file storage.

The communication channel requires a storage to persist messages that are relevant for a duration of time and it needs to support basic synchronization primitives to ensure consistency. Winslow instances that are started need to see events that were issued before their start and which are not completed yet to replicate the state of the system correctly.

For executing stages large files must be stored. The synchronized access is ensured by Winslow's event synchronization and thus the constraints to the storage is more relaxed than the one used to persist synchronization messages.

5.4.3   Layer 3: Backend Driver

This layer is responsible to interface with Docker. Once it is decided to run a stage this driver is instructed to start a certain image with environment variables and a prepared workspace. By separating this task to its own layer it will be easier to change from the Docker API to another implementation or away from the Docker platform at all if necessary. A reason for this could be the change of needs, the appearance of a platform that is fulfilling the needs better or the disappearance of the currently used platform. The latter sounds not that common at first, but the recent partial acquisition of Docker Inc by Mirantis[43] proves that even very popular third party software is not going to be around for all eternity. By moving the driver implementation into its own layer, but leaving the interface definition

in layer 1 allows the driver implementation to be swapped easily without the inner layer noticing. It then also complies to the "Dependency Inversion Principle".

While the backend driver does not access any storage itself, the stages need to read and write to the large file storage. Docker provides means to mount local directories or remote NFS shares as volumes into the container natively. The access from within the executed stage can then be guarded to limit read and write access to only the required working directories.

5.4.4   Layer 4: Client Communication

The final layer in this architecture is the client communication layer. This service is not crucial for the actual execution and may be disabled on Winslow instances that shall not response to a user request themselves. A reason for this could be, that SSL certificate for a sub-domain points to a specific Winslow instances and unsecure access is not desired. The REST API shall provide resources that can be called from the statical served Angular Web-Application[2] or further tools to upload or download files, to control stage execution and to monitor usage and logs.

---

[2]see section 3.4

5.5 Time Schedule

The following figure shows schedule planned for this work:
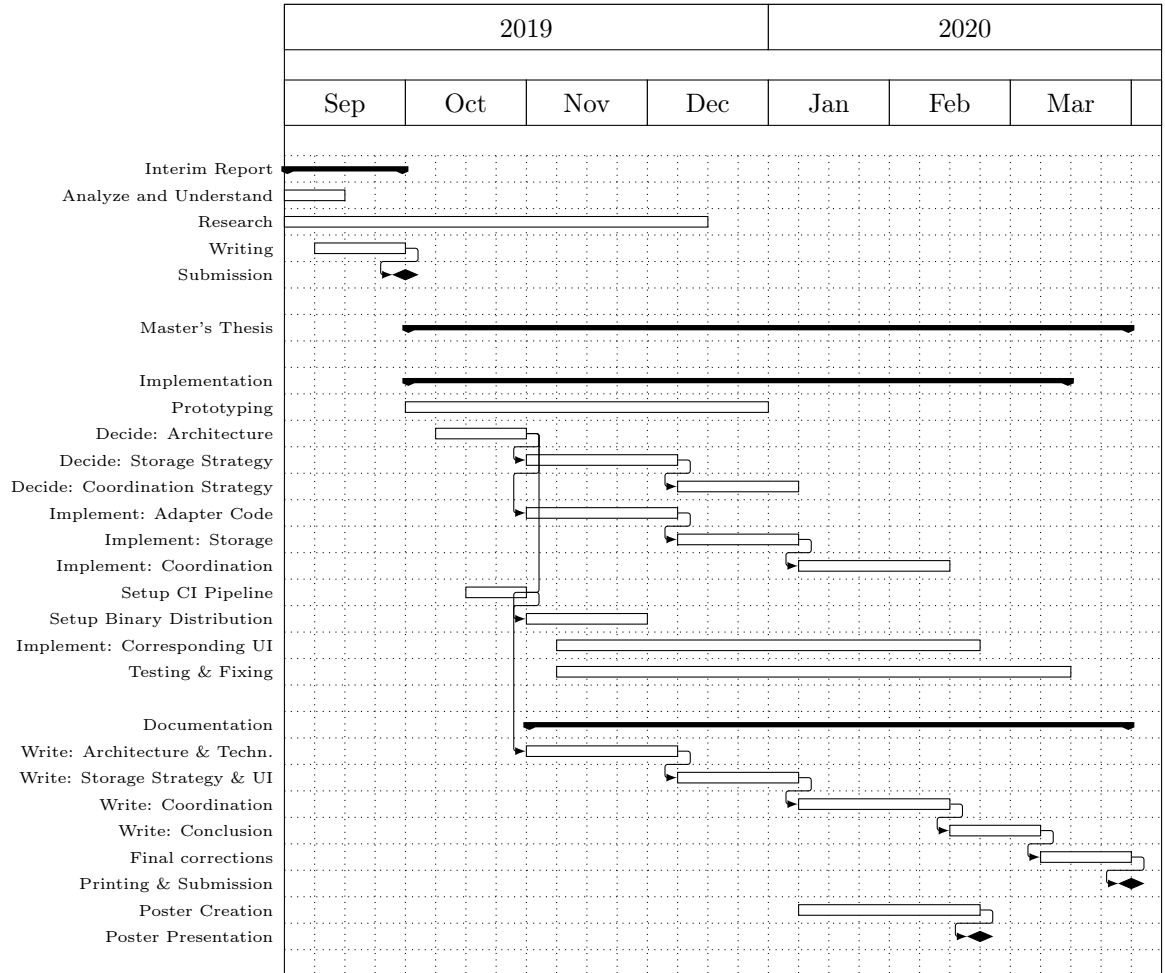
Figure 5.9: Time schedule

Chapter 6

System design

In this chapter the previously collected information are used to design the system in detail. The decision making process for that is documented here as well.

6.1 Environment

The environment of where the systems shall be deployed onto must be known for detailed design decisions, which will be elaborated here.

All Winslow instances are supposed to run inside a Docker container (as requested in section 2.2). The host will be a Ubuntu or Debian Linux Servers with some having multiple Nvidia graphics cards. There will be no graphical user interface available on these servers.

Furthermore, the decision to implement Winslow in Java was made early on. This has three main reasons, first, the developer is experience in Java development, the server side web framework used is a Java library (see section 6.10) and the company this implementation is for has a focus on Java development, which means if the system is successfully, maintained and possibly extended by another developer, it is highly probable that this developer has knowledge in Java development as well. TODO: .cleanup

6.2 Storage Technology

Storing and accessing data is one of the main concerns for Winslow. In almost all use cases, there will be multiple Winslow instances that need to access the storage. Because of that, it must be accessible from remote and from multiple Winslow instances simultaneously. For that, the following technologies are reviewed: NFS or SMB/CIFS (subsection 3.3.1), GlusterFS (subsection 3.1.6), SeaweedFS (subsection 3.1.6), OpenIO (subsection 3.1.6), dCache (subsection 3.1.5) and HadoopFS (subsection 3.1.1).

The following criteria were used for comparison, with a scoring system with points from 0 to 3, in which 0 means the best solution and 3 the worst - the lower the overall score the better:

- **Installation Overhead**: How hard and extensive is the installation? Is there a good documentation? Is the package provided from within the main repository of Ubuntu, or alternatively, does it provide any other form of easy installation? Scoring 0 for installation through the main repository, 1 for an external repository, 2 for distribution specific archive, 3 for any other installation.

- **Dependencies**: Are further services required for operation? How hard are they to setup? Do they need any additional configuration? Scoring 0 for no further dependencies, 1 for dependencies but without the need of configuration, 2 for dependencies that require manual configuration, 3 for dependencies that need to be installed manually and require manual configuration.

- **Single Point of Failure**: Is the solution decentralized and is it failure resilient? Is it site- or rack-aware or provide replication mechanisms to compensate? Scoring 0 for completely decentralized, 1 for decentralized but with dependency on centralized backend, 2 for mostly centralized but with load balancing approaches, 3 for no single point of failure mitigations.

- **Docker integration**: How easy is it to integrate with Docker? Scoring 0 for native support, 1 for native but non-trivial support, 2 for support through additional exports facility, 3 for non-trivial solution.

- **Failure concerns**: Are there any noteworthy concerns? Was an early local test successful and reliable? Is it a known technology or "proven in use"? Who is developing it and what are the support guarantees? Scoring 0 for commonly used and no concerns, 1 for minor concerns, 2 for major uncertainty, 3 for abandoned technology.

In Table 6.1 the results are displayed:

	NFS S.	SMB/CIFS S.	GlusterFS	SeaweedFs	dCache	HDFS
Inst.	0	1	0	2	3	1
Dep.	1	0	0	0	3	1
SPoF	2	3	0	2	0	0
Docker	0	2	2	3	2	3
Fail.	0	0	1	2	2	0
Score	3	6	3	9	10	5

Table 6.1: Comparison of storage technologies

The worst scoring tool in this comparison is dCache. The installation overhead, missing documentation and uncertainty on reliable operation is too high for this project (see subsection 3.1.5).

For similar reasons SeaweedFs is placed second worst. Local tests could not access SeaweedFs reliable when a node failed and there is no trivial support for Docker nor does it provide an NFS export. Using it from within Docker would require each container to be manipulated so that the first operation is mounting the storage with a custom binary and through a FUSE[1] mount. The tool also

---

[1]Filesystem in Userspace

seems to be developed by a single person which introduces further uncertainty about reliability and support in the future.

HDFS and SMB/CIFS seem to be no terrible choice, but no especially good one either in this use case.

The two best scoring storage solutions are a simple NFS share and GlusterFS. While GlusterFS provides replication and decentralized access, a NFS share scores with its simplicity. The idea is to start with a plain NFS share for Winslow which can natively be utilized by Docker as volume mount and to revisit later whether the need for replication and decentralization persists. Because of the NFS interface export of GlusterFS, Winslow could then easily switch to GlusterFS by accessing the NFS export closest to its' instance.

6.3   Execution Management

TODO: .polish

As noted in subsection 5.4.1, the central business logic is to deciding when and issuing stage executions. Generally speaking, there are two approaches when executing jobs: local or remote. Both will be discussed next.

In the remote approach, the job is executed on another machine and not on the same which has the responsibility to manage the process. Continuous Integration (CI) platform Jenkins[4] does offer this approach. The so called slave node (Jenkins) is accessed through a native interface (SSH for Linux servers) to copy resources and to start the job process. In this scenario, the CI instance requires and stores login credentials for every remote machine to be able to login whenever needed. The system administrator therefore has to create a new user account on the remote machine, install required programs and prepare the environments. One big risk for this approach is that in case of any security breach on the central CI instance, the attacker is also able to login on all remote machines.

GitLab[3] follows the approach, where the system administrator has to man-

ually install the GitLab Runner on the machine that is then able to connect to GitLab and execute jobs. This runner is responsible in pulling jobs, execute them locally, monitor and report back the outcome.

Winslow is going to follow the local approach. On each physical machine that is supposed to execute stages, Winslow needs to be installed and connected to the cluster. The installation is expected to be easy, as is based on starting prepared Docker images. An election is selecting the instance which is most fitting for a stage execution (see affinity and aversion in subsection 6.7.1). It is expected that each instance can judge this best on its own, as it knows which resources are available. Work is therefore distributed to an executor as it is detected.

6.4   Event Synchronization and Communication

As discovered in subsection 5.4.2 there is a need for an event synchronization across Winslow instances. Without proper coordination, race conditions could cause stages to be started multiple times simultaneously, corrupt workspaces, configuration or project files. While starting too many stage executions are only wasting resources, data corruptions can lead to unrecoverable damage.

There are multiple ways to exchange data between systems that do not share the same process or machine. The most flexible implementation can be achieved by implementing a custom protocol on a raw TCP socket which allows to exchange blobs (Binary Large Objects) between exactly two nodes. For a blob or message to reach all nodes, a connection to every other node, a centralized broker, another topology such as a tree structure or broadcast messages would be required. Because this in itself seemed to be a complex subject, third party alternatives were investigated.

Apache Kafka[44] is one such alternative. It is open source, distributed and is focused on providing system wide ordered data and event stream that can be persisted over a certain time period. But as every third party service, it intro-

duces a dependency on the project. Not only in regards to interfacing and driver implementation, but also for maintenance and setup. Each Winslow Image would require to be shipped with a pre-configured Kafka service and the administrator must ensure that these services can reach each other - in addition to storage solution.

TODO: .rework transition It would be rather elegant, if an existing connection - the storage - between all Winslow instances could be reused for this. This could eliminate the need for this third party dependency. To do so, a mechanism must be found to give each occurring event a sequence number, that is unique system wide. A new event would not be allowed to be propagated by an instance before it did not process all previous events. With this idea in mind, messages that can communication all required coordination information are defined next.

6.4.1 Messages

TODO: this chapter really needs further polishing

To coordinate the Winslow instances, messages will be exchanged through the common event bus. Because of the nature of this global event bus, messages are delivered by broadcasting them to all instances.

Because this is a multi instance system which can suffer partial failure, all multi-part operations have to have a timeout, so that a failure is detectable. Without this timeout detection, one could not detect the absence of finish signal of a multi-part operation, which could potentially block further operations for forever. The presence of a timeout will not be mentioned again when listing the messages in detail.

To account for transmission delay and clock offsets an additional time padding is granted. Because of the non-requirement of real-time scheduling this is generously set to 5 seconds.

A message has the following fields:

- **issuer**: The id of the Winslow instance that wrote the message

- **command**: The state change command (described a bit further below)

- **subject**: What the command is about to change

- **timestamp**: The time of when the message was issued

- **duration**: Optionally a non zero time period of the expected duration

With this message structure definition, the following commands can be used:

- **LOCK**: Locking a resource, commonly a project. A lock is exclusive to the issuer and can only be granted if there exist no lock for the same subject yet.

- **EXTEND**: Extending an ongoing operation by pushing back its timeout. This can be used to signal, that a lock is still required although its timeout is close The duration of the original event extended and thus the timeout is pushed back. The issuer must be the same as did issue the operation that is referred to.

- **RELEASE**: Releasing a locked resource. The issuer must be the same node as the one that issued the **LOCK** previously.

- **KILL**: Non-gracefully stop an operation or destruction of a lock. This signals that an operation shall be stopped or that the lock shall be released immediately. This can be used to abort an running stage execution.

- **ELECTION_START**: Signals that an election for a stage execution started. The signal must refer to an project that can make progress.

- **ELECTION_PARTICIPATE**: Signals that the issuer node is capable of executing the next stage of a project. It also includes a scoring for the affinity and aversion (more details in subsection 6.7.1). This message also includes the election it refers to.

- **ELECTION_STOP**: Signals that an election has finished. The participator with the best affinity and aversion scoring is now allowed execute the next stage of the referred project. This is signal is only allowed to be issued by the same node that started the election process.

6.4.2 Synchronization

In a file system two files with the same path are now allowed to exist, therefore the sequence number of the event to publish can be used as name for the file in a known directory to write events to. Watching the directory will then result a stream of ordered events. To publish an event, a naive implementation would fist check the directory for a file for the sequence number of the new event and create it, if it was not found. But this has a potential race condition: between checking for the file and creating it, another instance could have created the file as well. Luckily, the POSIX standard provides a flag when creating files, that will expose this conflict by returning an error if "Both O_CREAT and O_EXCL are set, and the named file already exists"[45]. This mechanism is also exposed in Java by calling `Files.write(..)` [46] with `StandardOpenOptions.CREATE_NEW` [47]. The behaviour of publishing an event can therefore be summarized in Figure 6.1:
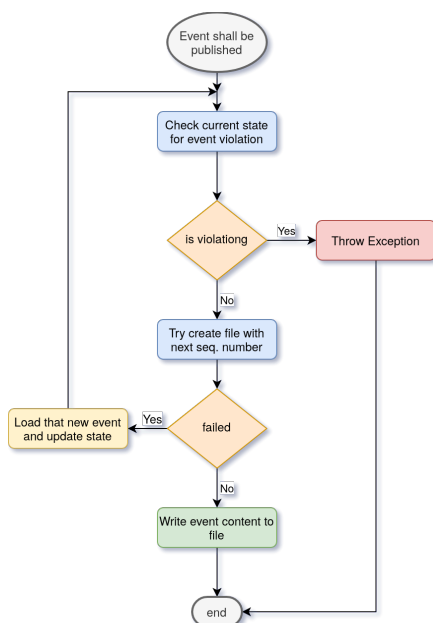
Every Winslow instance is responsible for keeping track of all global states and check before publishing a new event whether the state would be violated by this event. If it is not, it will try to publish it by creating a new file with the next expected sequence number. If the file already exists, another instance did publish an event in the meantime. This externally published event must then loaded and the global state updated before a new trial to publish its own event can be made. Eventually, this will either lead

Figure 6.1: Publishing process

to successfully publishing the event or failing due to state violation - for example trying to lock a project that is already locked.

This synchronization mechanism ensures that each event is only published if it is not violating any constraints, such as locking resources that are already locked. It also allows for multiple locks and elections to happen concurrently, as long as there is no overlap.


Figure 6.2: Sample event and lock series with lifetimes

In Figure 6.4.2 a series of events, their id and commands are shown. The coloured rectangles below shows the duration of the lifetime of events. The crossed-out event is violating the constraint that a resource can only be locked by one instance at a time, which is the owner of the yellow lifetime in this case. Because of this violation the even is never published to the event bus, which is demonstrate with the sequence number not skipping the number 4.

6.5 Docker interface

TODO: .cleanup

To interface with the Docker daemon a driver implementation is needed to pass on container creation. Docker provides a REST API[48] for third party libraries as well as implementations in Go and Python. There is no official Java API implementation. This means, either a third party Java library must be used, the REST calls that are required have to be implemented within this project or another alternative must be found. The Docker REST API is very expressive

which is another reason to find a viable alternative.

As it turns out, the Nomad Project, which was already investigated in subsection 3.1.4, provides a library with a thin abstraction layer and lots of reasonable default values. It also includes support to create container which use the Docker plugin by Nvidia[49]. This allows attaching GPUs to a container without the need of listing all libraries and device files manually[2]. The scheduling capabilities of Nomad are not used because the focus of nomad does only overlap partially with the needs of Winslow. A job is expected by Nomad to rerun, moved to different hosts and - in the service mode - scaled as it seems needed. The control of where the job is actually executed is to lost, which is not acceptable when different firewall rules apply to all Winslow instances. When using GlusterFS, other cloud storage or compute capabilities, the storage is not always reachable through the same path for every node. Nomad will only be used in the local Winslow container as thin layer to interfere with an enriched Docker API.

6.6 Directory Structure and Organization

This section describe the organization of the working directory for the Winslow instances. Because a common network share is required by Winslow for event synchronization and the projects' workspaces, this was further extended to share the configuration and project files as well. This has the very nice side-effect, that is no real setup to do when installing a new Winslow instance[3].

For the configuration files, a principle found on Unix and Linux systems was applied: human readable text files. For complex structures YAML formatting is used, while for simple key value files property formatting is used. This makes understanding state the system is for debugging in error scenarios easier.

---

[2]If you are interested in how tedious this could otherwise become see `https://github.com/hashicorp/nomad/issues/3499#issuecomment-364214506`

[3]Appendix A lists the complete installation script for a new Winslow instance

6.6.1 Winslow working directory

A brief summary of the by all instances shared working directory:

- `logs` : The location for stage log files. Logs system events and console output with timestamps. File name format looks like `<project-id>-<stage-number>-<stage-name>` .

- `pipelines` : Pipeline definition file (YAML) are located here.

- `projects` : Project definition files (YAML) are located here.

- `resources` : The globally available input resources, see subsection 6.6.2.

- `run/events` : The directory used to synchronized events, see section 6.4.

- `run/nodes` : The directory used to publish node utilization, see section 6.9.

- `settings` : The directory used to share common configurations, currently this only contains a global environment variable configuration.

- `workspaces` : Discussed in subsection 6.6.2

Because the files in `run/events` and especially `run/nodes` are very temporary, the NFS server locates these in shared memory (`/dev/shm`) to reduce stress on IO and unnecessary wear on SSDs. Thanks to the locking synchronization and locking mechanism (presented in section 6.4), the directory can shared between all Winslow instances.

6.6.2 Stage storage

Thinking about the storage organisation for the pipeline and its stages, a few expectations and concerns arise. First of all, to redo a stage, one needs to be able to access the files that were the result of one, two or multiple stages before the current. Sometimes a stage wants to access intermediate data produces by

multiple previous stages. Next, the input video footage needs to be accessed by multiple stages throughout the pipeline execution. Finally, some stage results are not intermediate but final results.

The first and second concern can be solved by providing a workspace directory for each stage, that is copied from the logically previous stage[4]. Once the computation of a stage is completed, the workspace is considered immutable and only used to source new workspaces from. This works fine for small intermediate results, but it does not work very well for large files - like the video footage. This delays the start of the stage execution, requires unnecessary storage due to multiple copies and provides no benefits in an archival and version control sense, because the video footage is not altered. So there needs to be another storage pool for input data, that is globally accessible and never changed: the global input storage pool. Providing one further storage pool for final results (global output pool), concern number three and four are also solved.

Because the very first stage has no workspace to source its files from, on creation of the pipeline a workspace directory for the "zeroths" stage is created. The user can then provide the very first stage with a predefined and non-empty workspace if necessary.

This also solves te problem with delays due copy operations in NFS, which are performed client side. This means, the client reads the input file and writes to the output file[5]. This operation does not only take unnecessary long but also utilizes all available network bandwidth which then cannot be used by other applications.

To ensure that the global input and the previous intermediate results are not altered, the Docker daemon is instructed to mounted them as read-only filesystems. This also prevents bugs or errors from accidentality deleting unrelated files[6]

---

[4]the stage the new one is based on, this does not always need to be the TODO: numerical previous stage

[5]Server-side copy has just been standardized for NFS 4.2[50]

[6]"This so theoretical and will actually never happen in real life" proven otherwise: `https://github.com/valvesoftware/steam-for-linux/issues/3671`

Figure 6.3: Stages with their intermediate and global resources

In summary, for a running stage the mounted directories look like the following::

- `/input` is mounted read-only from

  `nfs-server:/winslow/workspaces/<pipeline-id>/input`

- `/output` is mounted with write permissions from

  `nfs-server:/winslow/workspaces/<pipeline-id>/output`

- `/workspace` is mounted with write permissions from

  `nfs-server:/winslow/workspaces/<pipeline-id>/<stage-id>`

  and was created by the copying the workspace directory of the logically previous stage

6.7 Job Scheduling and Election System

Winslow has no ahead of time plan to schedule stage executions, the main reason for this that it is unknown how long a stage is about to take. One could consider to build statistics to eventually be able to guess the duration of repeatedly executed stage definitions, but this seems error prone, does not cover seldom executed stages and is in addition unreasonable complex. Instead, the system is reacting to an user submissions and on stages that have just completed. For each new

execution, an election is then held in which the winner will execute the stage. Winslow instances that do not fulfil the resource requirements do not participate in the election.

6.7.1 Affinity and Aversion

TODO: .

For a winner to be selectable, participants must be comparable. Winslow uses a two factor scoring mechanism for that judges how efficient a node is going to be utilized (affinity scoring) and which unused resources will be wasted (aversion scoring). The instance with the highest affinity out of the instances of the lowest aversion wins.

For determining the affinity scoring only requested resource categories c_i need to be considered. Through the resource monitoring (see section 6.9) the total resources $r_{c_{max}}$ as well as the reserved resources $r_{c_{in\_use}}$ is known. For every resource requested by the stage definition $c1..cx$, the ratio of request r_{c_i} to available resources is built. The lowest ratio, and therefore the most pessimistic value, of any category is used as the affinity score:

$$affinity = \min_{c_1..c_x} \left(\frac{r_{c_i}}{r_{c_{max}} - r_{c_{in\_use}}} \cdot n_{c_x} \right) \tag{6.1}$$

As seen above, every node can apply a node and resource category specific multiplier n_{c_x} for punishment or gratification. This allows to fine tune the score.

For determining the aversion scoring, all by a node available resource categories c_a need to be considered. In contrast to the affinity score, the highest ratio of any unused resource is used as aversion score.

$$aversion = \max_{c_1..c_a} \left(\frac{r_{c_{max}} - r_{c_{in\_use}} - r_{c_a}}{r_{c_{max}} - r_{c_{in\_use}}} \cdot n_{c_a} \right) \tag{6.2}$$

The more resources are wasted - especially categories that are untouched - the worse the aversion score.

An example shall illustrate the scoring system. Lets consider the following three execution nodes exist:

	Free CPUs	Free memory	Free GPUs
Node 1	4	16 GiB	0
Node 2	12	58 GiB	0
Node 3	4	50 GiB	1

Table 6.2: Nodes to consider

The following three stages executions shall be assessed:

	Req. CPUs	Req. memory	Req. GPUs
Ex1	2	16 GiB	0
Ex2	4	48 GiB	0
Ex3	4	16 GiB	1

Table 6.3: Stage executions to consider

The following table shows the resulting affinity and aversion score. The winning node for an execution is marked by the bold formatting.

	Node 1	Node 2	Node 3
Ex1	**0.5/0.5**	0.28/0.72	0.32/1.0
Ex2	-	**0.34/0.67**	0.96/1.0
Ex3	-	-	**0.32/0.68**

Table 6.4: Affinity/Aversion score for every stage definition and node constellation

As shown in the example, the aversion score is important to prevent stages taking execution nodes that provide specialised hardware (such as GPUs) but which isn't required by the stage execution. In this example, it prevent Ex2 to be scheduled on Node 3.

In comparison to just forbidding executions on nodes with hardware that is

not used, in scenarios where all general purpose nodes have failed, this system would still allow executions like Ex1 and Ex2 to be scheduled on GPU nodes.

6.8   Execution

The execution of a stage is started on the Winslow instance with the best aversion and affinity score and therefore won the election. It first locks the project, to re-read the configuration and to prevent any changes to in the meantime. The workspace is prepared (see subsection 6.6.2), environment variables collected and the backend driver issued to start the new container. Any data logged to the containers' stdout or stderr is collected and written to a log file exclusive to this stage execution (see subsection 6.6.1). System related events, such as errors pulling an image, exhausted resources or sudden aborts are also logged to this file.

The lock on the project configuration file is not held for all the time the stage is executed, because it would prevent the ability to enqueue and prepare further stages. Instead, after successfully starting the new container and the log file has been created, a new lock pointing to the log file is created. The currently executed stage is noted in the project file and then the lock on it is released. As long as the lock for the log file is alive, no election for the project is triggered.

Once the container stopped - and the stage has therefore either failed or succeeded - the lock on the project is acquired again to update the stage state. After this, both locks are released, potentially triggering the next stage execution election.

6.9   CPU, RAM, Network- and Disk-IO Monitoring

TODO: polish!

Monitoring the hardware utilization can be accomplished by parsing the contents of `/proc/cpuinfo` [7], `/proc/stat`, `/proc/meminfo` [8], `/proc/net/dev` and `/proc/diskstats` [9]. These special files in the `/proc` directory are text files generated by the Linux kernel to summarize the current utilization[51].

Each Winslow instance is repetitively reading and parsing these files and writing a summarized version to `/run/node/<node-name>` in common network share, an example can be seen in Appendix B.

6.10   User Interface

To display data[10] on the user interface in the Angular Web-Application, REST requests are sent to the web-backend of Winslow. To handle these requests and to send the responses SpringBoot[52] is used.

To not stress the event bus with lock requests caused by the user, all data retrieval operation open and read configuration files without locking them. This bears the risk of reading incompletely written files, but as the user interface is not crucial to the system it must account for failed requests due to network issues anyway, this is considered acceptable. The user interface would simple resubmit its request. TODO: ACID hazards ref?

When modifying settings or triggering stage executions, the request handling must also obey the locking rule to not introduce inconsistent state into the system (due to a dirty read). This means, on a HTTP POST request, the affected resource is being locked, updated and released. The release will then trigger a listening Winslow instance, which will check the file to potentially spawn an election.

---

[7]`https://www.kernel.org/doc/Documentation/cputopology.txt`
[8]`https://www.kernel.org/doc/Documentation/vm/hugetlbpage.txt`
[9]`https://www.kernel.org/doc/Documentation/ABI/testing/procfs-diskstats` and `https://www.kernel.org/doc/Documentation/iostats.txt`
[10]Screenshots in Appendix C

6.11   Continuous Deployment

Winslow uses GitLab CI to continuously deliver runnable Docker Images after each git push. Every code change will trigger a new build pipeline, that builds and tests Winslow, packages all dependencies[11] into a Docker image and pushes it to our in house private Docker Registry. Broken code or failing unit tests will result in a notification for the developer and prevents the image from being published.

To update an execution node, the Docker container can simple be stopped and discarded. The installation script[12] then pulls and starts the most recent image.

---

[11]Nomad and the Angular Web-Application

[12]Which is displayed in Appendix A

Chapter 7

Outcome and Measurements

In this chapter the outcome is measured and evaluated.

7.1 What did not work as expected

TODO: ?

It is impossible to plan a system ahead of time without making some mistakes. The biggest issue that surfaced on early implementation of Winslow were caused by too short-termed EXTEND signals (see subsection 6.4.1). In combination with short but high network bandwidth usages, the event was published too late, which led the other instances consider the node to have failed. This was solved by increasing lock timeouts and issuing the EXTEND signal halfway through the timeout instead of very close to the end.

7.2 User Authentication and Security

One aspect that was not yet mentioned at all is user authentication and security. Winslow is prepared in this regards because every project has an owner field and a member list and every web access check these. An internal user repository also associates user with groups and provides defaults, such as the user and group

"anonymous" and "root", for no and full access privileges, respectively. But at the moment, there is no user login, instead to the system all requests seem to be issued by "root". In the future this is planned to be replaced by an Single Sign-On implementation that uses the user accounts of our company. Winslow is then associating a from this service given user name to an internal user and groups. Unknown users are not allowed to see or alter anything, known users will be able to create and view their own projects and projects they are listed in as members.

Winslow also supports secure HTTP (HTTPS) whenever a SSL certificate is provided[1].

7.3   High Level System Overview

TODO: .delete?

Figure 7.1 shows the high level system overview of Winslow as it has been implemented by the end of this thesis in our company.

---

[1]This can be seen in Appendix A

Figure 7.1: High level view on the system structure

The user can access the system through the Web-Application, which itself is not indicating from which Winslow instance it is served. The three compute nodes supply different hardware, one with 128 GiB RAM, one with 4 NVIDIA Titans, and one with 48 Intel CPU cores. All instances access the same NFS share for synchronization and work distribution.

7.4 Failure resilience

TODO: .delete? or polish!

One of the initial requirements is to be resilient against node failures. Node failures can be simulated easily by killing the Docker container of a Winslow instance. The behaviour that is then displayed depends on whether that killed instance was actually executing a stage. If it did not, the only change is that in the Web-Application the node with its utilization reports will vanish.

But if it did execute a stage, it held locks for it. These locks will expire,

and after the additional grace period (see subsection 6.4.1), the other instances will notice it, lock the affected project themselves and mark the execution as failed. Currently, the project is then paused and a user confirmation awaited. An alternative to this could be to re-schedule the stage automatically.

7.5 Evaluation

The main focus of Winslow is to improve the efficiency in utilizing available hardware. Without it, new stages were seldom started at outside of the work time. To estimate whether Winslow is providing any improvements, the actions that are scheduled by Winslow outside of these hours are to be watched.

The following graphs are based on activity metrics, that were logged in the time period from the first usage in mid January 2019 until counter actions for COVID-19 in early March 2019 prevented active use due to the lock-down in Germany. For each day of the week and hour of the day an associated cell shows the activity colour encoded. Furthermore, red lines show the begin and end of a typical work day. Interestingly, the actual position of these red lines can also be inferred by the graphs themselves, due to activity spikes.

Figure 7.2: Actively running stages at a given day of the week and time of day

At first, Figure 7.2 and Figure 7.3 show the overall activity. In Figure 7.2 number of running stages per hour is shown and in Figure 7.3 the number of started stages per hour is shown. The increased activity between the read line - within the working hours - is visible. But especially in Figure 7.3 the activity

outside these times is visible as well.

Figure 7.3: Stages starting at a given day of the week and time of day

Figure 7.4 and Figure 7.5 show only stages that were triggered by Winslow, without a direct user input.


Figure 7.4: Actively running stages at a given day of the week and time of day which were started automatically

Figure 7.2 clearly shows that there is quite a bit of activity event without direct user input. This is especially obvious by looking at Figure 7.5: there are many stages started in the late evening and until midnight as well as a second wave at about one o'clock in the morning and activity on Saturdays. Without Winslow these stages would not have been triggered.

That the system was actively used to schedule work to be executed on the weekends. This can be seen in Figure 7.5 by highlighted areas late hours on Fridays, early and late hours on Saturdays as well as the continuous but decreasing overall activity from Saturday to Sunday in Figure 7.4. The gap of activity on Mondays at around noon in Figure 7.4 could indicate that the results were evaluated at here.


Figure 7.5: Stages automatically starting at a given day of the week and time of day

For the graphs above, the table can be used to summarize the results:

Overall in the test period 2003 hours were spent computing - this includes overlapping stage executions - of which 78.6% were executed within regular working hours. To see the benefits of Winslow, the hours spent in stage executions that

	absolute	relative	gain
All	2003h	100%	
Within work time	1575h	78.6%	
Outside work time	428h	21.4%	
Triggered by Winslow overall	1131h	56.5%	+129%
Triggered by Winslow outside work time	286h	14.3%	+16%

Table 7.1: Categorization of stage execution hours

were automatically issued are of interest. Here, Winslow triggered 1575 hours or 78.6% of all stage executions which did not require staff assistance. 286 hours were overall gained by automatically triggering and performing stage executions outside of the working hours of the staff. This is an increase of +16% compared to no automated triggers.

To put this into perspective, typically week consists of 40 hours working and 128 hours with the staff being absent from the office. So in theory 128 hours (+220%) per week could additionally be spent executing stages, which raises the question why the numbers of Winslow are so low.

To answer this, one needs to look again at the heat-maps of Figure 7.2 and Figure 7.4. Winslow is being used, and as for Figure 7.5 discussed, it is being used for long running stage executions with one or two automatic triggers at nights. But it is used even more during working hours, as can be seen in Figure 7.2, where all pink cells fall within working hours.

In conclusion, the Winslow is being used - actively, with over 2003 hours worth of compute time[2]. It is supporting the users by automatically triggering over half of all compute time and additionally enabling the staff to let large multi-stage tasks be executed over night or the weekend. If the system would not be beneficial to the staff, it would not be used in such an extend so shortly after implementation.

---

[2]in the previously mentioned time period

Chapter 8

Further work

In the scope of this thesis a system was implemented to distribute the workload of a computer vision pipeline onto several compute nodes. Automatically stage scheduling improves the utilization of available hardware - especially at times where the staff is absent, like at nights or weekends.

Adding the support for WebSockets could help to make the Web-Application more responsive to serve side changes. Pushing desktop notifications to the client or e-mail messages could reduce the response time of the staff for when a stage failed or a project needs attention otherwise. GlusterFs could be investigated for whether the (site) replication feature could reduce the usage of network bandwidth. A system like this provides an endless stream of ideas to further improve the user experience or for extensions and capabilities to add.

Chapter 9

Conclusion

Overall the project was a success. Winslow helps to improve the usage of available hardware resources by additionally triggering stage executions in times where the staff is absent or not watching. It helps in keeping order by separating workspaces and organizing them into projects. The user interface provides an interaction method that is easy to understand for users, and provides upload and download mechanism without resorting back to console commands.

In the implementation phase an interesting and minimalistic way to synchronize events based on a commonly shared directory was discovered and implemented. Winslow is easy to setup and due to its few dependencies and easy to maintain. The evaluation showed that Winslow is increasing productivity and hardware utilization without additional attention by the staff and that it is beneficial to the user.

Bibliography

[1] MEC-View Dr. Rüdiger W. Henn. *Mobile Edge Computing basierte Objek-terkennung für hoch- und vollautomatisiertes Fahren.* URL: `http://mec-view.de/` (visited on 09/29/2019).

[2] The Apache Software Foundation. *Apache Hadoop.* URL: `https://hadoop.apache.org/` (visited on 09/19/2019).

[3] Inc. GitLab. *The first single application for the entire DevOps lifecycle.* URL: `https://about.gitlab.com/` (visited on 09/21/2019).

[4] jenkins.io. *Jenkins. Build great things at any scale.* URL: `https://jenkins.io/` (visited on 09/19/2019).

[5] Inc. GitLab. *GitLab CI/CD. Pipeline Configuration Reference.* URL: `https://docs.gitlab.com/ee/ci/yaml/` (visited on 09/19/2019).

[6] jenkins.io. *Using a Jenkinsfile.* URL: `https://jenkins.io/doc/book/pipeline/jenkinsfile/` (visited on 09/19/2019).

[7] Camunda Services GmbH. *Workflow and Decision Automation Platform.* URL: `https://camunda.com/` (visited on 09/19/2019).

[8] Camunda Services GmbH. *Process Engine API.* URL: `https://docs.camunda.org/manual/7.6/user-guide/process-engine/process-engine-api` (visited on 09/19/2019).

[9] Camunda Services GmbH. *Rest Api Reference.* URL: `https://docs.camunda.org/manual/7.8/reference/rest` (visited on 09/19/2019).

[10] HashiCorp. *Nomad.* URL: `https://www.nomadproject.io/` (visited on 09/19/2019).

[11] HashiCorp. *Nomad. Nomad vs. Kubernetes.* URL: `https://www.nomadproject.io/intro/vs/kubernetes.html` (visited on 09/19/2019).

[12] Bc. Pavel Peroutka. *Web interface for the deployment andmonitoring of Nomad jobs. Master's thesis.* URL: `https://dspace.cvut.cz/bitstream/handle/10467/80106/F8-DP-2019-Peroutka-Pavel-thesis.pdf` (visited on 09/22/2019).

[13] Tigran Mkrtchyan Patrick Fuhrmann. *dCache. Scope of the project.* URL: `https://www.dcache.org` (visited on 09/19/2019).

[14] Patrick Fuhrmann. *dCache, the Overview.* URL: `https://www.dcache.org/manuals/dcache-whitepaper-light.pdf` (visited on 09/19/2019).

[15] Inc. Terracotta. *QuartzJob Scheduler.* URL: `http://www.quartz-scheduler.org/` (visited on 09/19/2019).

[16] Data Revenue. *Distributed Python Machine Learning Pipelines.* URL: `https://www.datarevenue.com/en/blog/how-to-scale-your-machine-learning-pipeline` (visited on 09/19/2019).

[17] Ask Solem. *Celery: Distributed Task Queue.* URL: `http://www.celeryproject.org` (visited on 09/19/2019).

[18] IBM Knowledge Center. *IBM InfoSphere. DataStage.* URL: `https://www.ibm.com/support/knowledgecenter/en/SSZJPZ_9.1.0/com.ibm.swg.im.iis.ds.design.doc/topics/c_ddesref_Server_Job_Stages_.html` (visited on 09/19/2019).

[19] Wikipedia contributors. *Qsub. Wikipedia, The Free Encyclopedia.* URL: `https://en.wikipedia.org/w/index.php?title=Qsub&oldid=745279355` (visited on 09/22/2019).

[20] The University Of Iowa. *Basic Job Submission. HPC Documentation Home / Cluster Systems Documentation.* URL: `https://wiki.uiowa.edu/display/hpcdocs/Basic+Job+Submission` (visited on 09/22/2019).

[21] Swiss National Supercomputing Centre. *High Throughput Scheduler.* URL: `https://user.cscs.ch/tools/high_throughput/` (visited on 09/19/2019).

[22] Colin Phipps. *zsync. Overview.* URL: `http://zsync.moria.org.uk/` (visited on 09/19/2019).

[23] OpenIO. *High Performance Object Storage for Big Data and AI.* URL: `https://www.openio.io/` (visited on 09/22/2019).

[24] Chris Lu. *Simple and highly scalable distributed file system.* URL: `https://github.com/chrislusf/seaweedfs` (visited on 09/22/2019).

[25] Colin Phipps. *Alluxio. Data Orchestration for the Cloud.* URL: `https://www.alluxio.io` (visited on 09/19/2019).

[26] Inc Red Hat. *Gluster. Free and open source scalable network filesystem.* URL: `https://www.gluster.org` (visited on 09/19/2019).

[27] Inc. Docker. *Enterprise Container Platform.* URL: `https://www.docker.com/` (visited on 09/22/2019).

[28] Inc. Docker. *Docker. Docker Logos and Photos.* URL: `https://www.docker.com/company/newsroom/media-resources` (visited on 01/16/2020).

[29] Inc. Docker. *Docker Documentation. Swarm mode key concepts.* URL: `https://docs.docker.com/engine/swarm/key-concepts/` (visited on 03/17/2020).

[30] The Kubernetes Authors. *Kubernetes. What is Kubernetes.* URL: `https://kubernetes.io/docs/concepts/overview/what-is-kubernetes/` (visited on 03/17/2020).

[31] Inc. Docker. *Docker Hub. Build and Ship any Application Anywhere.* URL: `https://hub.docker.com/` (visited on 09/22/2019).

[32] Neil Brown. *Linux Kernel Documentation. Overlay Filesystem.* URL: `https://www.kernel.org/doc/Documentation/filesystems/overlayfs.txt` (visited on 03/29/2020).

[33] Inc. Docker. *Docker Documentation. Docker overview.* URL: `https://docs.docker.com/engine/docker-overview/` (visited on 01/16/2020).

[34] IETF | Internet Engineering Task Force. *RFC 5661. Network File System (NFS) Version 4 Minor Version 1 Protocol.* URL: `https://tools.ietf.org/html/rfc5661%5C#section-1.7.2.2` (visited on 03/26/2020).

[35] Google. *Angular.* URL: `https://angular.io/` (visited on 03/29/2020).

[36] Microsoft. *TypeScript. JavaScript that scales.* URL: `https://www.typescriptlang.org/` (visited on 03/29/2020).

[37] IETF | Internet Engineering Task Force. *RFC 7231. Hypertext Transfer Protocol (HTTP/1.1).* URL: `https://tools.ietf.org/html/rfc7231#section-4` (visited on 03/29/2020).

[38] J. Goll. *Entwurfsprinzipien und Konstruktionskonzepte der Softwaretechnik: Strategien für schwach gekoppelte, korrekte und stabile Software.* Springer Fachmedien Wiesbaden, 2018.

[39] R.C. Martin, J.M. Rabaey, A.P. Chandrakasan, et al. *Agile Software Development: Principles, Patterns, and Practices.* Alan Apt series. Pearson Education, 2003.

[40] J. Goll. *Methoden des Software Engineering: Funktions-, daten-, objekt- und aspektorientiert entwickeln.* Springer Fachmedien Wiesbaden, 2012.

[41] "Use Case Modeling". In: *Use Case Driven Object Modeling with UML: Theory and Practice.* Berkeley, CA: Apress, 2007, pp. 49–82.

[42] Jeffrey Palermo. *The Onion Architecture.* URL: `https://jeffreypalermo.com/2008/07/the-onion-architecture-part-1/` (visited on 03/26/2020).

[43]   Inc. Docker. *Press Release. Docker Restructures and Secures $35 Million to Advance Developer Workflows for Modern Applications*. URL: `https://www.docker.com/press-release/docker-new-direction` (visited on 03/23/2020).

[44]   The Apache Software Foundation. *Apache Kafka*. URL: `https://kafka.apache.org/` (visited on 03/29/2020).

[45]   GNU.org / The Free Software Foundation. *Opening and Closing Files*. URL: `https://www.gnu.org/software/libc/manual/html_node/Opening-and-Closing-Files.html` (visited on 03/29/2020).

[46]   Oracle. *Java Platform SE 7. Files*. URL: `https://docs.oracle.com/javase/7/docs/api/java/nio/file/Files.html` (visited on 03/29/2020).

[47]   Oracle. *Java Platform SE 7. StandardOpenOptions*. URL: `https://docs.oracle.com/javase/7/docs/api/java/nio/file/StandardOpenOption.html#CREATE_NEW` (visited on 03/29/2020).

[48]   Inc. Docker. *Develop with Docker Engine API*. URL: `https://docs.docker.com/engine/api/` (visited on 03/29/2020).

[49]   NVIDIA. *NVIDIA Container Toolkit*. URL: `https://github.com/NVIDIA/nvidia-docker` (visited on 03/29/2020).

[50]   IETF | Internet Engineering Task Force. *RFC 7862. Network File System (NFS) Version 4 Minor Version 2 Protocol*. URL: `https://tools.ietf.org/html/rfc7862` (visited on 03/29/2020).

[51]   *The /proc filesystem*. URL: `https://www.kernel.org/doc/Documentation/filesystems/proc.txt` (visited on 03/29/2020).

[52]   Inc. or its affiliates VMware. *Spring*. URL: `https://spring.io/` (visited on 03/29/2020).

Appendix A

Winslow Instance Installation

At the time of writing this, the following script is all that is required to install a new Winslow instances:

```bash
#!/bin/bash

NODE_TYPE="executor"
STORAGE_TYPE="nfs"
STORAGE_PATH="johnny5.itd-intern.de:/data/streets/winslow"
ADDITIONAL="-p 446:4646 -e WINSLOW_DEV_ENV=true -e
    WINSLOW_DEV_REMOTE_USER=root"

PARAMS="$@"
IMAGE="repo.itd-intern.de/winslow/node"
NODE_NAME="$(hostname)"
CONTAINER_NAME="winslow"
GPUS="$(ls /dev/ | grep -i nvidia | wc -l)"
WORKDIR="/winslow/"

SUDO=""

if [ "$(id -u)" -ne 0 ] && [ "$(id --name -G | grep -i docker | wc
    -l)" -eq 0 ]; then
    SUDO="sudo"
```

```
19  fi
20
21  $SUDO docker pull $IMAGE
22
23  if [ "$KEYSTORE_PATH_PKCS12" != "" ]; then
24      ADDITIONAL="$ADDITIONAL -p $HTTPS:8080 -v
        $KEYSTORE_PATH_PKCS12:/keystore.p12:ro -e
        SERVER_SSL_KEY_STORE_TYPE=PKCS12 -e
        SERVER_SSL_KEY_STORE=file:/keystore.p12 -e
        SECURITY_REQUIRE_SSL=true -e SERVER_SSL_KEY_STORE_PASSWORD="
25  fi
26
27
28  echo ""
29  echo ""
30  echo ""
31  echo " :::::   Going to create Winslow Container with the following
        settings"
32  echo ""
33  echo "   HTTP Port     '$HTTP'"
34  echo "   HTTPS Port    '$HTTPS'"
35  echo "    Docker Image '$IMAGE'"
36  echo "    Storage Type '$STORAGE_TYPE' @ '$STORAGE_PATH'"
37  echo ""
38  echo "    Work Directory '$WORKDIR'"
39  echo "    Node Name      '$NODE_NAME'"
40  echo ""
41  echo "    Detected GPUs: $GPUS"
42  echo ""
43  echo "    Additional Docker Parameters:   '$ADDITIONAL'"
44  echo "    Additional Winslow Parameters: '$PARAMS'"
45  echo ""
46
47  sleep 1
48
```

```
49  if [ $($SUDO docker ps --filter "name=$CONTAINER_NAME" | wc -l) -gt
      1 ]; then
50         echo " ::::: Stopping already running Winslow instance"
51             $SUDO docker stop "$CONTAINER_NAME" > /dev/null && echo
      " :::: Done" || (echo " :::: Failed"; exit 1)
52      fi
53
54  $SUDO docker rm "$CONTAINER_NAME" > /dev/null
55
56  echo " ::::: Starting Winslow Container now"
57  $SUDO docker run -itd --privileged \
58      --restart=unless-stopped \
59      --name "$CONTAINER_NAME" \
60      $(if [ "$GPUS" -gt 0 ]; then echo "--gpus all"; fi) \
61      -p $HTTP:8080 \
62      $ADDITIONAL \
63      -e WINSLOW_STORAGE_TYPE=$STORAGE_TYPE \
64      -e WINSLOW_STORAGE_PATH=$STORAGE_PATH \
65      -e WINSLOW_WORK_DIRECTORY=$WORKDIR \
66      -e "WINSLOW_NODE_NAME=$NODE_NAME" \
67      $(if [ "$NODE_TYPE" == "observer" ]; then echo "-e
      WINSLOW_NO_STAGE_EXECUTION=1"; fi) \
68      -v /var/run/docker.sock:/var/run/docker.sock \
69      $IMAGE \
70      $PARAMS
71
72  $SUDO docker attach "$CONTAINER_NAME"
```

Listing A.1: The whole installation script for Winslow

Appendix B

Node Status Information

?? shows what a status file of an execution node in `/run/nodes/` looks like:

```
1  name = "srv515"
2
3  [cpuInfo]
4  modelName = "Intel(R) Core(TM) i7-6850K CPU @ 3.60GHz"
5  utilization = [0.029411765, 0.009803922, 0.0, 0.009803922, 0.0,
       0.0, 0.0, 0.0, 0.0, 0.029411765, 0.0, 0.0]
6
7  [memInfo]
8  memoryTotal = 67447889920
9  memoryFree = 55159222272
10 systemCache = 39576674304
11 swapTotal = 1074786304
12 swapFree = 815677440
13
14 [netInfo]
15 transmitting = 2888
16 receiving = 2042
17
18 [diskInfo]
19 reading = 0
20 writing = 188416
```

```
21  free = 440000196608
22  used = 48637153280
23
24  [[gpuInfo]]
25  vendor = "nvidia"
26  name = "GeForce RTX 2080 Ti"
27
28  [[gpuInfo]]
29  vendor = "nvidia"
30  name = "GeForce RTX 2080 Ti"
31
32  [[gpuInfo]]
33  vendor = "nvidia"
34  name = "GeForce RTX 2080 Ti"
35
36  [[gpuInfo]]
37  vendor = "nvidia"
38  name = "GeForce RTX 2080 Ti"
39
40  [buildInfo]
41  date = "2020-02-12 14:24:39"
42  commitHashShort = "cb259423"
43  commitHashLong = "cb259423a47a38530abae211cf0108fbaf01d852"
```

Listing B.1: Sample for a node status information

Appendix C

Screenshots

Figure C.1: Left: Known projects, Right: Overview for the selected project
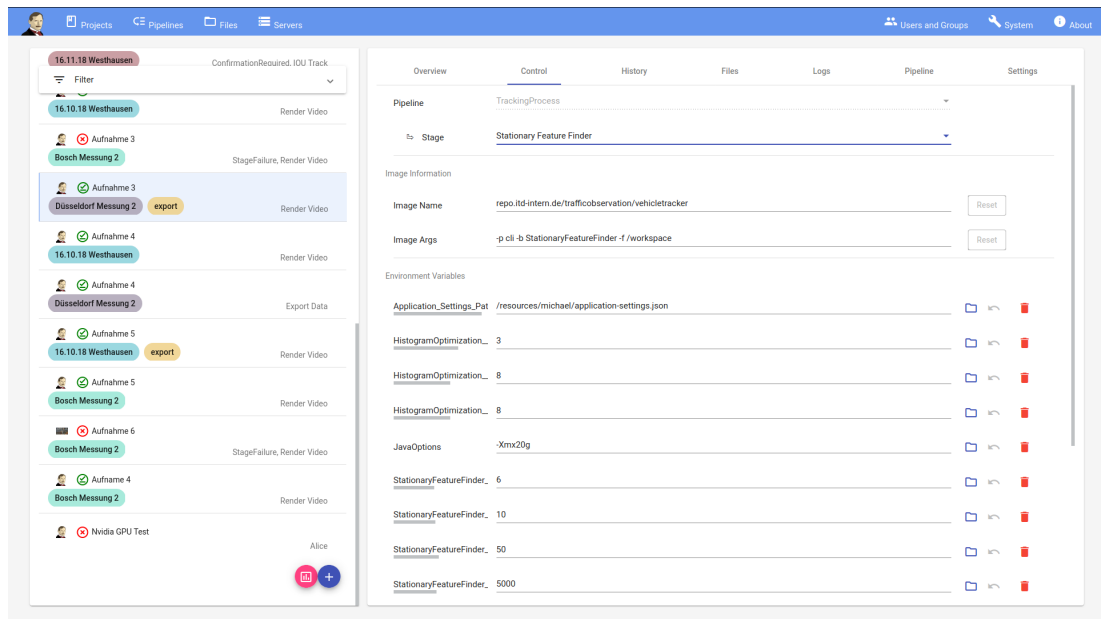
Figure C.2: Left: Known projects, Right: Control for the selected project

Figure C.3: Left: Known projects, Right: History of the selected project

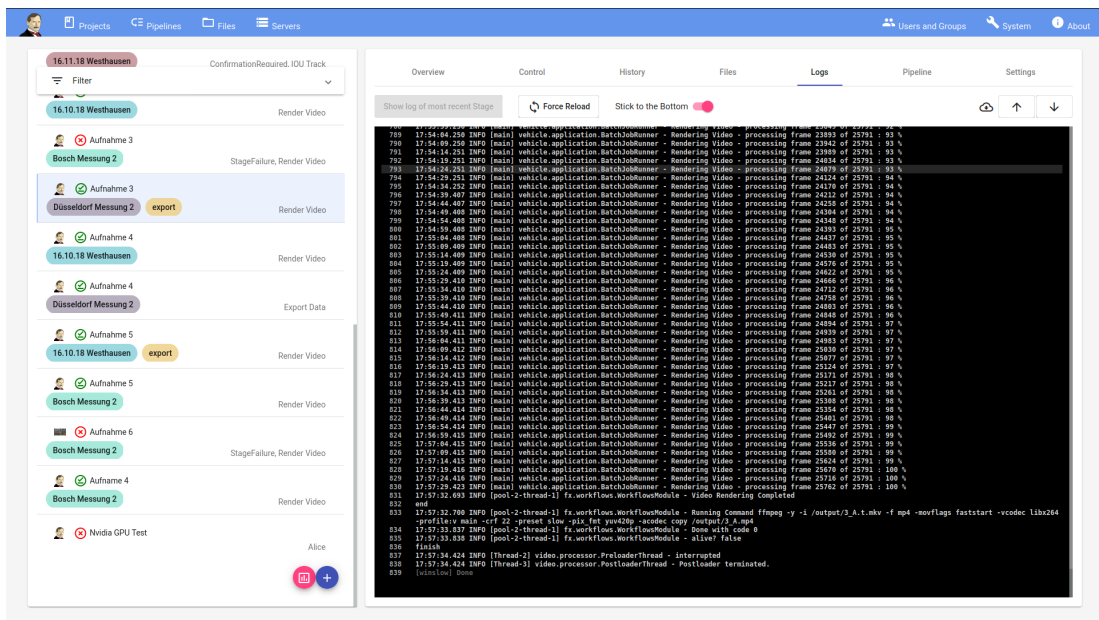Figure C.4: Left: Known projects, Right: Files associated with the selected project

Figure C.5: Left: Known projects, Right: Most recent logs for the selected project

Figure C.6: Left: Known projects, Right: Pipeline of the selected project
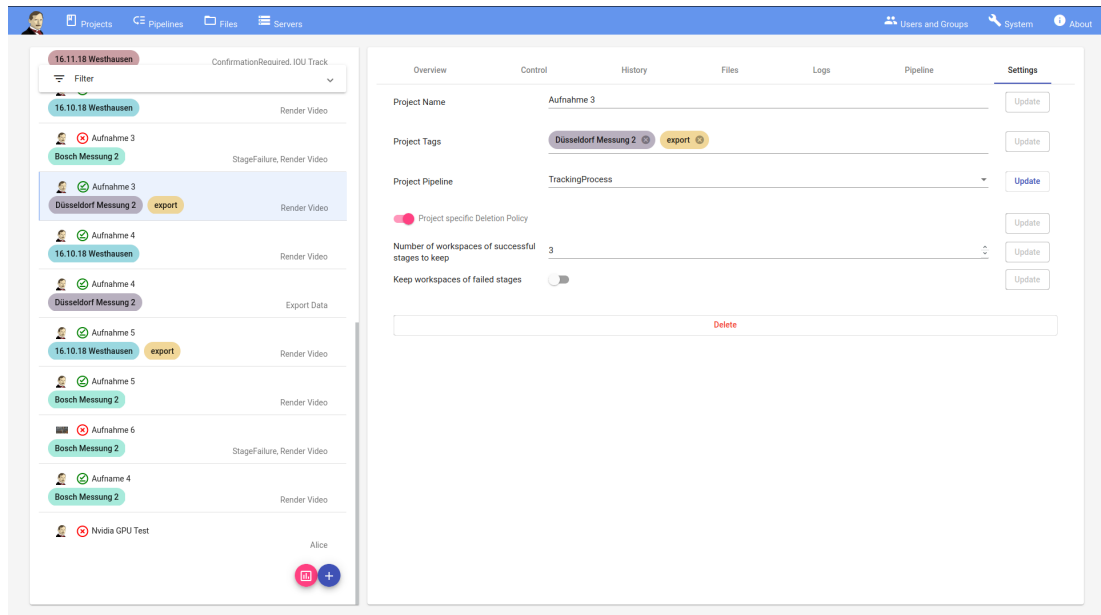

Figure C.7: Left: Known projects, Right: Settings of the selected project

Figure C.8: Edit view for pipeline definitions

Figure C.9: File explorer which is not restricted to a project

Figure C.10: Utilization overview for every attached Winslow instance

Figure C.11: Utilization overview for every attached Winslow instance

Appendix D

Interim Report

Contents

Chapter 1

Introduction

Since the industrial revolution, humans strive for more automation in the industry as well as in the every day life. What was at first a cost saving measurement in factories, now also is a differentiation method for products. A new product must prove a higher level comfort to the customer than the previous generation as well as all the competitors. As such, the ambitions of the industry are focused on increasing the value of their products for the customer.

The automotive industry is one of the prime examples of this. Never was traveling from one place to another as comfortable as nowadays. Aspects like an elegant interior design, comfortable seats, air conditioning, entertainment systems and safety measurements need to be considered by car manufacturers to be competitive these days. The next luxury enhancement will be the autonomously driving vehicle. No longer shall the owner of a car steer it, but instead the car becomes his or hers personal chauffeur, driving the optimal route, the most comfortable way and being more reliable and safer than any human ever could.

The reason, autonomously driving cars are not common already, is their big complexity increase. Compared to already established technologies like parking assistants, entertainment systems or more efficient engine controllers, letting a computer reliably understand a certain traffic situation requires masses of input data and complex algorithms to process. As such, the problem itself becomes massive and cannot be solved that easily. To solve this, the industry has no choice than to divide this into many small pieces and work out solutions to it step by step.

The MEC-View research project explorers one such step: whether and how to include external, steady mounted sensors in the decision finding process for partially autonomous vehicles in situations where onboard sensors are insufficient. To not disrupt traffic flow with non-human behavior, one needs to study and thereby watch human traffic. Automatically analyzing traffic from video footage requires a lot of computation power and can be further optimized by specialized hardware such as GPUs.

This thesis will conceptualize and realize a distributed and automated computer vision pipeline which can be used to analyzes traffic flow within video footage. Compared to an existing but highly manual workflow, the new system shall help to utilize the available hardware more efficiently by reducing idle times. Stage transitions and basic scheduling shall be automated to allow a user to plan and execute multiple projects ahead of time and in parallel.

1.1 MEC-View

The MEC-View research project[1] - funded by the German Federal Ministry for Economic Affairs and Energy - aims to supplement the field of view of automated driving cars with road-side sensor data using 5G mobile communication. The sensor information is merged into an environment model on the so-called Mobile Edge Computing (MEC) server. This server is directly attached to the radio station to ensure low latency environment model updates.

The project is tested at an intersection in Ulm, Germany. Currently, there are 15 lidar and video sensors installed. Those sensors send their detections to the (MEC) server. A fusion-algorithm merges those detections into one environment model and sends it back to the (MEC) server and to the automated cars.

Additionally, general traffic flow is analyzed to learn about movement patterns. To do so, 4k video data is captured by an air drone from real world cross roads. On each frame of such a recording, cars are detected with a neuronal network. Detected cars are tracked throughout the video to compute the movement speed and position in time of each car. In an analysis of all vehicles, hot-spots of high and low traffic flow can be determined.

Chapter 2

The Program

This chapter will discuss the program which shall be implemented. To do so, the problem to solve must be understood. To gather requirements and understand the technical hurdles to overcome, this chapter is split into two sections. First, a rough glance over the current workflow is given, which is followed by a more detailed description for the desired workflow.

2.1   Current Workflow

Currently, to analyze a video for the trajectories of the recorded vehicles, the following steps are executed manually:

1. Upload the input video to a new directory on the GPU server

2. Execute a shell script with the video as input file and let it run (hours to days) until completed. The shell script invokes a Java Program - called TrackerApplication - with parameters on what to do with the input file and additional parameters.

3. The intermediate result with raw detection results is downloaded to the local machine and opened for inspection. If the detection error is too high, the camera tracking has a drift or other disruptions are visible, the previous step is redone with adjusted parameters.

4. Upload the video and intermediate result to a generic computing server and run data cleanup and analysis. This is achieved with the same Java

Program as in step 2, but with different stage environment parameters.

5. Download the results, recheck for consistency or obvious abnormalities. Depending on the result, redo step 2 or 4 with adjusted parameters again.

6. Depending on the assignment, steps 4 and 5 are repeated to incrementally accumulate all output data (such as statistics, diagrams and so on).

Because all those steps are done manually, the user needs to check for errors by oneself. Also, if a execution is finished or failed early, there could be hours wasted if the regular check intervals are too far apart, such as during nights.

2.2   Desired Workflow

The desired workflow shall be supported through a rich user interface. This user interface shall provide an overview of all active projects and their current state, such as running computation, awaiting user input, failed or succeeded.

To create a new project, a predefined pipeline definition shall be selected as well as a name chosen. Because only a handful of different pipeline definitions are expected, the creation of such does not need to happen through the user interface. Instead, it is acceptable to have to manually edit a configuration file in such rare circumstances.

Once a project is created, the user wants to select the path to the input video. This file has to be been uploaded to a global resource pool at this point. The upload and download of files shall therefore also be possible through the user interface. Because a video is usually recorded in 4k (3840 x 2160 pixels), encoded with H.264 and up to 20 minutes long, the upload must be capable of handling files which are tens of gigabytes large.

Once a pipeline is started, it shall execute the stages on the most fitting server node until finished, failed or a user input is required. Throughout, the logs of the current and previous stage shall be accessible as well as uploading or downloading files from the current or previous stages workspace. In addition to the pipeline pausing itself for user input, the user shall be able to request the pipeline to pause after the current stage at any moment. When resuming the pipeline, the user

might want to overwrite the starting point to, for example, redo the latest stage.

Mechanisms for fault tolerance shall detect unexpected program errors or failures of server nodes. Server nodes shall be easily installed and added to the existing network of server nodes. Each server node might provide additional hardware (such as GPUs), which shall be detected and provided.

For the ease of installation and binary distribution, Docker Images shall be used for running the Java Program for analyzing the videos as well the to be implemented management software.

2.3 Deliverable Requirements

From the desired workflow, the following requirements can be extracted (shortened and incomplete due to early project stage):

- Rich user interface

- Storage management for global resource files as well as stage based workspaces

- Pipeline definition through configuration files

- Handling of multiple projects with independent progress and environment

- Reflecting the correct project state (running, failed, succeeded, paused)

- Log accumulation and archiving

- Accepting user input to update environment variables, resuming and pausing projects as well as uploading and downloading files into or from the global resource pool or a stages workspace.

- Assigning starting stages to the most fitting server node

- Detecting program errors (in a stage execution)

- Cope with server node failures

- Docker Image creation for the Java Binary as well as the program implementation, preferred in an automated fashion.

Chapter 3

State of the art

In this chapter, programs solving similar problems, as described in the desired workflow, or dealing with a subset of the problem are looked into. The reason for this is to use well established or suitable programs as middle-ware to reduce implementation overhead. Where this is not possible, one might be able to gather ideas and learn about proven strategies to use or pitfalls to avoid while implementing custom solutions.

3.1 Similar solutions

This sections focuses on programs trying to provide somewhat similar workflows.

3.1.1 Hadoop MapReduce

For big data transformation, Hadoop MapReduce[2] is well known. With MapReduce, the input data is split into blocks of data and distributed onto Mappers. Mappers then apply the business logic and output intermediate results in form of key/value pairs. After shuffling, the Reduce stage will combine values from key/value pairs with the same key. In the final output, each key is unique accompanied with a value.

This strategy has proven to be very powerful to process large amount input data because Mappers and Reducers can work independently on their data-sets and therefore scale very well when adding further instances.

If the implementation were to be based on Hadoop MapReduce to achieve the desired workflow, it could be done like the following:

- Each video is split into many frames and each frame is applied to a Mapper

- A Mapper tries to detect all vehicles on a frame and outputs their position, orientation, size and so on

- The Reducer then tries to link the detections of a vehicle through multiple frames

- The final result would be a set of detections and therefore all positions for each vehicle in the video

But at the moment, this approach seems to be unfitting due to at least the following reasons:

1. It is not always trivial to reasonable link the detections of a vehicle. For example, a vehicle can be hidden behind a tree for a few frames until visible again. In addition, MapReduce requires the combination to be performed per common key. Until one is trying to link the detections of multiple frames, there is no common identifier that could easily be used as key. The position of a moving vehicle cannot be used as key, neither can the color or size, because of the noise of the camera, deviation in detection output and perspective distortions. The current implementation of the TrackerApplication is archiving this by finding similarities between detections, but for the Mapper it would be required to express this as a deterministic key.

2. MapReduce is great in combining many machines to solve a big computational problem. But at the moment, this is neither a desired nor given condition. At the moment, there are a handful of very powerful workstations with specialized hardware. Therefore it is perfectly acceptable and sometimes required, when each workstation works through a complete video at a given time instead.

3.1.2 Build Pipelines

Build pipelines such as GitLab[3] and Jenkins[4] can also distribute the execution of stages onto other server nodes. In a common use-case, such build pipelines are used to build binaries out of source code, after a new commit into a SCM[1] repository was made. At IT-Designers GmbH GitLab as well as Jenkins are commonly used for scenarios exactly like this. A pipeline definition in GitLab CI/CD [5] or in a Jenkinsfile [6] describe stages and commands to execute. Each stage can be hosted on another node and be executed sequential or in parallel to each other.

Although this seems to be quite fitting for the desired workflow, there are two issues. First of all, such a pipeline does not involve any user input besides an optional manual start invocation. The result is then determined based on the state of the input repository. Second, such a pipeline is designed to determine the output (usually by compiling) whereas each run is independent from the previous and a repeated run shall provide the same result as the previous did. Usually, a new run is only caused by a change of the input data. However, the desired workflow differs in this aspects. A redo of a stage can depend on the result of the previous stage, for example, if the results are poor or the the stage failed. Instead of having multiple complete pipeline runs per project, the desired workflow uses a pipeline definition as base for which the order can be changed. Also, intermediate results need to influence further stages, even if repeated.

3.1.3 Camunda

Camunda[7] calls itself a "Rich Business Process Management tool" and allows the user to easily create new pipelines by combining existing tasks with many triggers and custom transitions. Camunda is focused upon visualizing the flow and tracking the data through a pipeline. The Camundas Process Engine[8] also allows user intervention between tasks.

One of the main supporting reason for it Camunda is the out of the box rich graphical user interface for process definition and interaction. Through its API[9], Camunda also allows custom external workers to execute a task. But it misses the

---

[1]Source Code Management

capability to control which task shall be processed on which worker node which is required by the desired workflow. It does also not provide any concept on how to allocate and distribute resources. The user interface - while being rich overall - is quite rudimentary when it is about configuring tasks and would therefore require custom plugins to be developed for more advanced user interactions.

Camunda is also not designed to reorder stages or insert user interactions at seemingly random fashion. The user itself is considered more as a worker that gets some request, "executes" this externally and finally marks the request as accepted or declined. Mapping this to the desired workflow does not feel intuitive. Finally, there is also no overview of task executors, no centralized log accumulation and no file up- or download for global project resources.

3.1.4 Nomad

Nomad[10] by HashiCorp is a tool to deploy, manage and monitor containers, whereas each job is executed in its own container. It provides a rich REST API and can consider hardware constraints on job submissions. Compared to Kubernetes[11], which is similar but more focused on scaling containers to an externally applied load, it is very lightweight. It is also available in many Linux software repositories - such as for Debian - which makes the installation very easy.

Because there were no grave disadvantages found (depending on a third party library can always be considered be a disadvantage for flexibility, error-pronous and limit functionality) Nomad is being considered as a middle-ware to manage and deploy stages. Others[12] seem to be using Nomad to manage and deploy containers for similar reasons. Nonetheless, further testing and prototyping will be required for a final decision.

3.1.5 dCache

"The goal of this project is to provide a system for storing and retrieving huge amounts of data, distributed among a large number of heterogenous server nodes, under a single virtual filesystem tree with a variety of standard access methods"[13]. dCache seems to be able to solve the storage access and distribution concern for the stages and sever nodes. When using dCache, one could store the

global resources distributed between the server nodes. Built-in replication would prevent access loss on a node or network failure and an export through NFS[2] allows easy access for Linux based systems[14].

But the installation is complex and requires many services to be setup correctly, such as postgresql and many internal services such as zookeeper, admin, poolmanager, spacemanager, pnfsmanager, cleaner, gplazma, pinmanager, topo, info and nfs. The documentation is also rather outdated and incomplete which meant, early tests with a prototype setup took days to setup and behaved rather unstable (probably due to a wrong configuration). It is to be seen, whether such an complex and heavy system is actually required or if there are feasible alternatives.

3.1.6 Further mentions

The following list shall acknowledge programs that behave similar to the previously mentioned strategies. Programs that are listed here, were looked into, but not in-depth because miss-fits were detected early on (listed in no specific order):

- **Quartz**[15] is a Java based program to schedule jobs. Instead of doing so by using input, Quartz executes programs through a timetable and in certain intervals.

- **Luigi**[16] also executes pipelines with stages and is written in python. The advertised advantage is to define the pipeline directly in python code. But, this is at the same time the only way to define pipelines which contradictions with the existing Java TrackerApplication implementation.

- **Calery**[17] is focused on task execution through message passing and is written in Python. Intermediate results are expected to be transmitted through messages. Because there is no storage strategy and python adapter-code would have been required, Calery was dismissed.

- **IBM InfoSphere**[18] provides similar to Camunda a rich graphical user interface but for data transfer. Dismissed due to commercial nature.

---

[2]Network File System

- **qsub**[19][20] is a CLI[3] used in HPC to submit jobs onto a cluster or grid. Dismissed due to an expected high setup overhead, non-required multi-user nature and the fact, that it only provides a way to submit jobs.

- **CSCS**[21] High Throughput Scheduler (GREASY). Dismissed for similar reasons as qsub, although it is more light weight and hardware agnostic (it can consider CUDA/GPU requirements).

- **zsync**[22], similar to rsync, is a file transfer program. Zsync allows to only transfer new parts when a file that shall be copied already exists in an older version on the target. This tool might be useful when implementing a custom resource distribution strategy is required.

- **OpenIO**[23] provides a distributed file system, is already provided as Docker image and provides a simple to use CLI. Because the NFS export is only available through a paid subscription plan, it was dismissed from further investigation.

- **SeaweedFS**[24] provides a scalable and distributed file system. The most interesting aspects are that it is rack-aware as well as natively supports external storage such as Amazon S3. When adding server nodes from the cloud this could allow all nodes to access the same file system while using rack-aware replication to reduce bandwidth usage and latency. A local test also proved that it is easy to setup, but because it cannot hot-swap nodes and was not able to recover when the seeweed master node became unreachable it was dismissed.

- **Alluxio**[25] provides a distributed file system but was dismissed because it itself requires a centralized file system for the master and its fallback instances

- **GlusterFS**[26] is another tool to provide a distributed file system with replication. It was bought by IBM but is nonetheless available through the software repository of many Linux distributions such as Debian. A local test showed that the setup is very easy and no adjustments of configuration files

---

[3]Command Line Interface

are required. However, the replication mechanism requires that an integer multiple of nodes of the replica value are assigned to the file system. This makes GlusterFS hard to use in a scenario, where adding and removing nodes are expected to happen frequently. It was therefore dismissed.

3.2   Docker Integration

As describe before (see section 2.2), for easy deployment, the implementation as well as the stages shall be executed inside Docker[27] containers. This allows easier isolation of the stages and workspaces from each other and other host programs. Because one needs to communicate with the Docker daemon, this increases the complexity for the implementation. But by using third party libraries, the increase in complexity can be limited.

Chapter 4

Outcome and further work

In this chapter the main concerns are listed. For each concern the current progress is described as well as further work that needs to be done.

4.1 Storage

One of the central concerns is the storage management. The program needs to make input files available on each execution node and collect the results once the computation is complete. There are a few main architectural strategies to approach this. Simplified, either at a centralized location which is accessed by all execution nodes, a copy of the input files to the execution nodes or decentralized and distributed between all execution nodes and replication. The advantages and disadvantages can depend on the specific implementation and is therefore discussed in combination of such (see chapter 3).

Further testing is required to decide whether a more complex storage system is required, or the simplicity of a centralized solution outweighs the setup and maintenance overhead.

4.2 Coordination

Another important concern is the coordination of the nodes. A central coordinator with external server nodes, such as GitLab and Jenkins have, might not be sufficient for more complex and longer lasting pipelines. The probability that the

master would need to be offline while there is a stage executed, is in the scenario of the desired workflow higher than for GitLab or Jenkins, because the stage is being execution for hours or days. Coupling stage execution plans on node availability ahead of time, as well as recovering from a sudden master failure implies additional implementation complexity. A decentralized coordination needs to be able to do this as well, but also allows the usage of the system while a node failed or is unreachable due to maintenance. With further prototyping and research a reasonable solution shall be found.

4.3 Binary distribution

In a time where containers are common and have proven to be usable, the installation of the binaries directly on the operating system they are executed on shall be avoided. There shall be no manual, nor automatic but custom file copies of the binaries or images from one server to the other. Experience shows, that without a proper management, this can easily become a mess, in which it is no longer clear, which files or images belong to which version. At the same time, making all binaries publicly available through the Docker Hub[28] is no option either. Whether a self hosted Docker Registry[29] could be the solution to this will be determined in further testing.

4.4 User Interface

Providing a useful user interface might not be important to the functionality of the system itself but for the user experience. A bad user experience will cause a system not to be used. It became common practice for a rich user experience to be web based and interactive with JavaScript. For a potentially decentralized system, it is also advantageous to be able to access a disconnected node in the same manner as the remaining system, which further encourages a web based solution. Web based solutions such as React and Angular shall therefore be investigated for being used as user interface.

Chapter 5

Schedule

The following figure shows schedule for further and past work:

Figure 5.1: Time schedule

Bibliography

[1] MEC-View Dr. Rüdiger W. Henn. *Mobile Edge Computing basierte Objekterkennung für hoch- und vollautomatisiertes Fahren.* URL: `http://mec-view.de/` (visited on 09/29/2019).

[2] The Apache Software Foundation. *Apache Hadoop.* URL: `https://hadoop.apache.org/` (visited on 09/19/2019).

[3] Inc. GitLab. *The first single application for the entire DevOps lifecycle.* URL: `https://about.gitlab.com/` (visited on 09/21/2019).

[4] jenkins.io. *Jenkins. Build great things at any scale.* URL: `https://jenkins.io/` (visited on 09/19/2019).

[5] Inc. GitLab. *GitLab CI/CD. Pipeline Configuration Reference.* URL: `https://docs.gitlab.com/ee/ci/yaml/` (visited on 09/19/2019).

[6] jenkins.io. *Using a Jenkinsfile.* URL: `https://jenkins.io/doc/book/pipeline/jenkinsfile/` (visited on 09/19/2019).

[7] Camunda Services GmbH. *Workflow and Decision Automation Platform.* URL: `https://camunda.com/` (visited on 09/19/2019).

[8] Camunda Services GmbH. *Process Engine API.* URL: `https://docs.camunda.org/manual/7.6/user-guide/process-engine/process-engine-api` (visited on 09/19/2019).

[9] Camunda Services GmbH. *Rest Api Reference.* URL: `https://docs.camunda.org/manual/7.8/reference/rest` (visited on 09/19/2019).

[10] HashiCorp. *Nomad.* URL: `https://www.nomadproject.io/` (visited on 09/19/2019).

[11] HashiCorp. *Nomad. Nomad vs. Kubernetes*. URL: `https://www.nomadproject.io/intro/vs/kubernetes.html` (visited on 09/19/2019).

[12] Bc. Pavel Peroutka. *Web interface for the deployment andmonitoring of Nomad jobs. Master's thesis*. URL: `https://dspace.cvut.cz/bitstream/handle/10467/80106/F8-DP-2019-Peroutka-Pavel-thesis.pdf` (visited on 09/22/2019).

[13] Tigran Mkrtchyan Patrick Fuhrmann. *dCache. Scope of the project*. URL: `https://www.dcache.org` (visited on 09/19/2019).

[14] Patrick Fuhrmann. *dCache, the Overview*. URL: `https://www.dcache.org/manuals/dcache-whitepaper-light.pdf` (visited on 09/19/2019).

[15] Inc. Terracotta. *QuartzJob Scheduler*. URL: `http://www.quartz-scheduler.org/` (visited on 09/19/2019).

[16] Data Revenue. *Distributed Python Machine Learning Pipelines*. URL: `https://www.datarevenue.com/en/blog/how-to-scale-your-machine-learning-pipeline` (visited on 09/19/2019).

[17] Ask Solem. *Celery: Distributed Task Queue*. URL: `http://www.celeryproject.org` (visited on 09/19/2019).

[18] IBM Knowledge Center. *IBM InfoSphere. DataStage*. URL: `https://www.ibm.com/support/knowledgecenter/en/SSZJPZ_9.1.0/com.ibm.swg.im.iis.ds.design.doc/topics/c_ddesref_Server_Job_Stages_.html` (visited on 09/19/2019).

[19] Wikipedia contributors. *Qsub. Wikipedia, The Free Encyclopedia*. URL: `https://en.wikipedia.org/w/index.php?title=Qsub&oldid=745279355` (visited on 09/22/2019).

[20] The University Of Iowa. *Basic Job Submission. HPC Documentation Home / Cluster Systems Documentation*. URL: `https://wiki.uiowa.edu/display/hpcdocs/Basic+Job+Submission` (visited on 09/22/2019).

[21] Swiss National Supercomputing Centre. *High Throughput Scheduler*. URL: `https://user.cscs.ch/tools/high_throughput/` (visited on 09/19/2019).

[22] Colin Phipps. *zsync. Overview*. URL: `http://zsync.moria.org.uk/` (visited on 09/19/2019).

[23] OpenIO. *High Performance Object Storage for Big Data and AI*. URL: `https://www.openio.io/` (visited on 09/22/2019).

[24] Chris Lu. *Simple and highly scalable distributed file system*. URL: `https://github.com/chrislusf/seaweedfs` (visited on 09/22/2019).

[25] Colin Phipps. *Alluxio. Data Orchestration for the Cloud*. URL: `https://www.alluxio.io` (visited on 09/19/2019).

[26] Inc Red Hat. *Gluster. Free and open source scalable network filesystem*. URL: `https://www.gluster.org` (visited on 09/19/2019).

[27] Inc. Docker. *Enterprise Container Platform*. URL: `https://www.docker.com/` (visited on 09/22/2019).

[28] Inc. Docker. *Docker Hub. Build and Ship any Application Anywhere*. URL: `https://hub.docker.com/` (visited on 09/22/2019).

[29] Inc. Docker. *Docker Documentation. Docker Registry*. URL: `https://docs.docker.com/registry/` (visited on 09/22/2019).